

Constraint satisfaction problem (CSP)

- Formalization of a CSP
- Backtracking algorithm
- Forward checking algorithm
- GAC algorithm

Next: Knowledge representation and reasoning (KR)

Readings: Chapter 8, 9

*Slides based on those of Sheila McIlraith

Constraint Satisfaction Problems

- The search algorithms we discussed so far had no knowledge of the states representation (black box).
- Instead we can have a general state representation that works well for many different problems.
- We can build then specialized search algorithms that operate efficiently on this general state representation.
- We call the class of problems that can be represented with this specialized representation CSPs – Constraint Satisfaction Problems.

CSPs: State Representation

The idea: represent states as vectors of feature values

- A set of k features (or variables)
- Each variable has a domain of different values, e.g.
 - height = {short, average, tall},
 - weight = {light, average, heavy}
- A state is specified by an assignment of a value for each variable.
- A partial state is specified by an assignment of a value to some of the variables.
- a goal state specified as conditions on the vector of feature values.

Example: Sudoku

	2							
		6				3		
7	4		8					
			3			2		
8		4			1			
6		5						
		1		7	8			
5			9					
				4				

1	2	6	4	3	7	9	5	8
8	9	5	6	2	1	4	7	3
3	7	4	9	8	5	1	2	6
4	5	7	1	9	3	8	6	2
9	8	3	2	4	6	5	1	7
6	1	2	5	7	8	3	9	4
2	6	9	3	1	4	7	8	5
5	4	8	7	6	9	2	3	1
7	3	1	8	5	2	6	4	9

Example: Sudoku

- 81 variables, each representing the value of a cell.
- Domain of Values: For cells that are already filled in, the domain is the singleton set containing the predefined value. The domain of other variables is the set {1-9}.
- State: any completed board given by specifying the value in each cell (1-9, or blank).
- Partial State: some incomplete filling out of the board.
- Solution: a value for each cell satisfying the constraints:
 - no cell in the same column can have the same value.
 - no cell in the same row can have the same value.
 - no cell in the same sub-square can have the same value.

Example: Scheduling

Want to schedule a time and a space for each final exam so that

- No student is scheduled to take more than one final at the same time.
- The space allocated has to be available at the time set.
- The space has to be large enough to accommodate all of the students taking the exam.

Example: Scheduling

We use the index i to uniquely identify exams – the “ i -th exam”.

- T_i is a variable representing the scheduled time for the i -th final exam.
 - Domains of T_i are fixed to {FriAm, TuesPm, MonAm, MonPm, …, FriPm}.
- S_i is the space variable for the i -th final.
 - Domain of S_i are all rooms big enough to hold the i -th final
- For all pairs of finals i, j ($i \neq j$) such that there is a student taking both: $T_i \neq T_j$
- For all pairs of finals i, j ($i \neq j$): $T_i \neq T_j$ or $S_i \neq S_j$

Formalization of a CSP

- A CSP consists of
 - A set of variables V_1, \dots, V_n
 - For each variable a (finite) domain of possible values $\text{Dom}[V_i]$.
 - A set of constraints C_1, \dots, C_m .
- A solution to a CSP is an assignment of a value to all of the variables such that every constraint is satisfied.
- A CSP is unsatisfiable if no solution exists.

Formalization of a CSP

- Each variable can be assigned any value from its domain
 - $V_i = d$ where $d \in Dom[V_i]$
- Each constraint C
 - has a set of variables it is over, called its scope, e.g., $C(V_1, V_2, V_4)$
 - Acts as a boolean function that maps assignments to these variables to true/false, e.g.,
 - $C(V_1 = a, V_2 = b, V_4 = c) = True$
 - $C(V_1 = b, V_2 = c, V_4 = c) = False$

Constraints

- Unary Constraints (over one variable)
 - e.g., $C(X) : X = 2$; $C(Y) : Y > 5$
- Binary Constraints (over two variables)
 - e.g., $C(X, Y) : X + Y < 6$
- Higher-order (n -ary) constraints: over 3 or more variables.

Specification of constraints

We can specify the constraint with a table

E.g., $C(V1, V2, V4)$ with

- $\text{Dom}[V1] = \{1, 2, 3\}$ and
- $\text{Dom}[V2] = \text{Dom}[V4] = \{1, 2\}$

V1	V2	V4	$C(V1, V2, V4)$
1	1	1	False
1	1	2	False
1	2	1	False
1	2	2	False
2	1	1	True
2	1	2	False
2	2	1	False
2	2	2	False
3	1	1	False
3	1	2	True
3	2	1	True
3	2	2	False

Specification of constraints

Often we can specify the constraint more compactly with an expression:

E.g., $C(V1, V2, V4) = (V1 = V2+V4)$

V1	V2	V4	C(V1,V2,V4)
1	1	1	False
1	1	2	False
1	2	1	False
1	2	2	False
2	1	1	True
2	1	2	False
2	2	1	False
2	2	2	False
3	1	1	False
3	1	2	True
3	2	1	True
3	2	2	False

Example: Sudoku

- **Variables:** $V_{11}, V_{12}, \dots, V_{21}, V_{22}, \dots, V_{91}, \dots, V_{99}$
- **Domains:**
 - $\text{Dom}[V_{ij}] = \{1-9\}$ for empty cells
 - $\text{Dom}[V_{ij}] = \{k\}$ a fixed value k for filled cells.
- **Constraints:**
 - Row constraints:
 - All-Diff($V_{11}, V_{12}, V_{13}, \dots, V_{19}$)
 - All-Diff($V_{21}, V_{22}, V_{23}, \dots, V_{29}$)
 -, All-Diff($V_{91}, V_{92}, \dots, V_{99}$)
 - Column Constraints:
 - All-Diff($V_{11}, V_{21}, V_{31}, \dots, V_{91}$)
 - All-Diff($V_{12}, V_{22}, V_{32}, \dots, V_{92}$)
 -, All-Diff($V_{19}, V_{29}, \dots, V_{99}$)
 - Sub-Square Constraints:
 - All-Diff($V_{11}, V_{12}, V_{13}, V_{21}, V_{22}, V_{23}, V_{31}, V_{32}, V_{33}$)
 - All-Diff($V_{14}, V_{15}, V_{16}, \dots, V_{34}, V_{35}, V_{36}$)

8		4	6		7
				4	
	1			6	5
5	9		3	7	8
				7	
4	8		2	1	3
	5	2			9
		1			
3		9	2		5

Solving CSPs

- We do not care about the sequence of moves needed to get to a goal state.
- We only care about finding a setting of the variables that satisfies the goal.
- Thus CSPs can be solved by a specialized version of depth first search.

Solving CSPs

- We can build up to a solution by searching through the space of partial assignments.
- In principle, the order in which we assign the variables does not matter – eventually they all have to be assigned.
- If we falsify a constraint during the process of building up a solution, we can immediately reject the current partial assignment.

CSP as a Search Problem

- Initial state: empty assignment
- Successor function: a value is assigned to any unassigned variable, which does not cause any constraint to return false.
- Goal test: the assignment is complete

A generic backtracking algorithm

- We pick a variable*,
- pick a value for it*,
- test the constraints that we can,
- if a constraint is unsatisfied we backtrack,
- otherwise we set another variable.
- When all the variables are set, we're done.

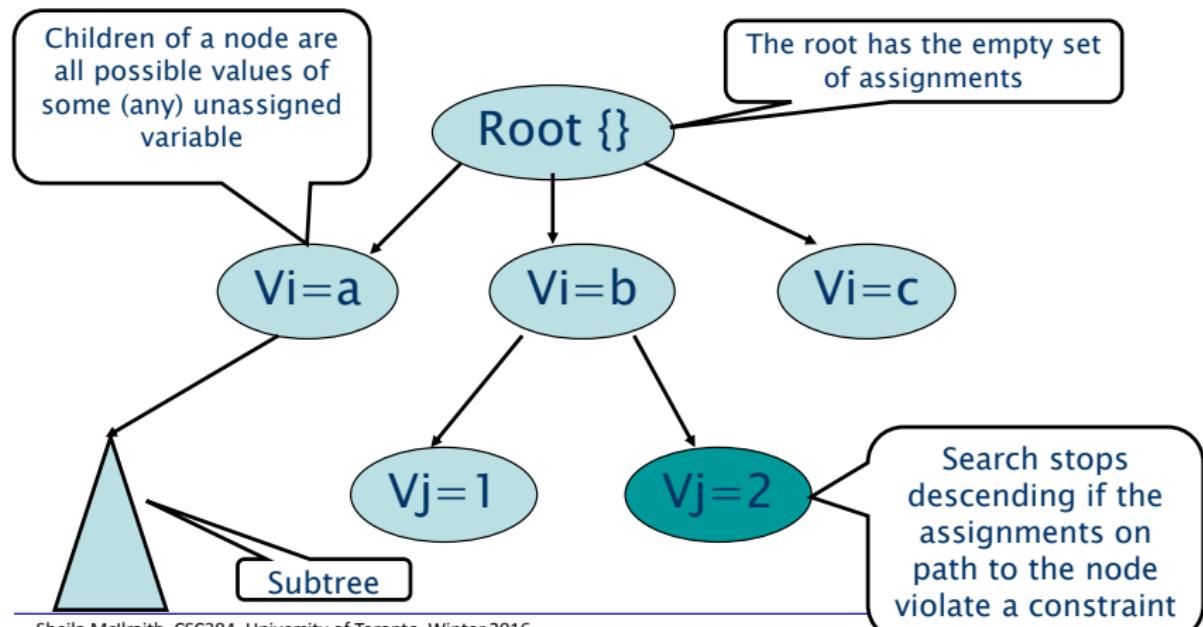
Backtracking Search: The Algorithm BT

```
BT(Level)
  If all variables assigned
    PRINT Value of each Variable
    RETURN or EXIT (RETURN for more solutions)
      (EXIT for only one solution)
  V := PickUnassignedVariable()
  Assigned[V] := TRUE
  for d := each member of Domain(V) (the domain values of V)
    Value[V] := d
    ConstraintsOK = TRUE
    for each constraint C such that
      a) V is a variable of C and
      b) all other variables of C are assigned:
          ;(rarely the case initially high in the search tree)
      IF C is not satisfied by the set of current
        assignments:
        ConstraintsOK = FALSE
    If ConstraintsOk == TRUE:
      BT(Level+1)

  Assigned[V] := FALSE //UNDO as we have tried all of V's values
  return
```

Backtracking Search

- The algorithm searches a tree of partial assignments.

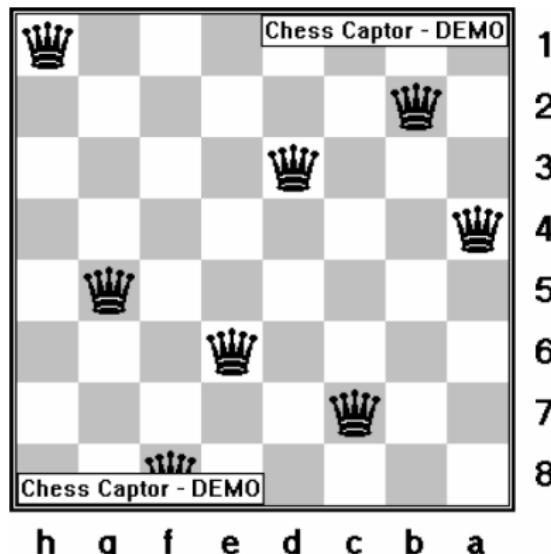


Backtracking Search

- Heuristics are used to determine
 - the order in which variables are assigned:
PickUnassignedVariable()
 - the order of values tried for each variable.
- The choice of the next variable can vary from branch to branch, e.g.,
 - under the assignment $V1=a$ we might choose to assign $V4$ next, while under $V1=b$ we might choose to assign $V5$ next.
- This “**dynamically**” chosen variable ordering has a tremendous impact on performance.

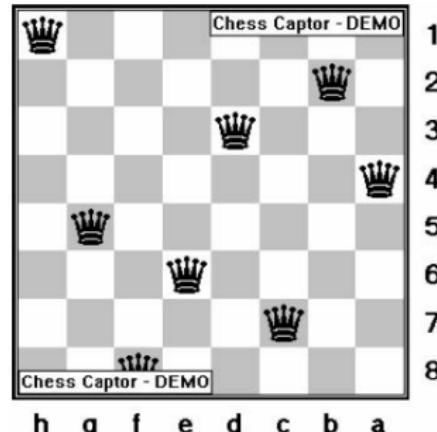
Example: N-Queens

Place N Queens on an N X N chess board so that no Queen can attack any other Queen.



Formulation 1

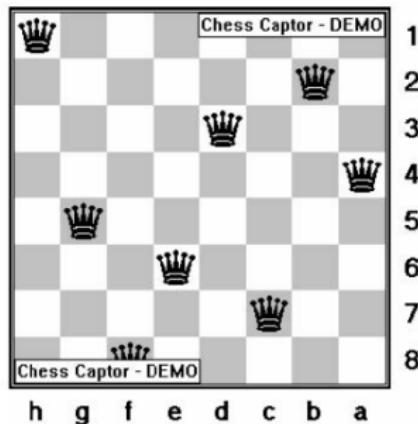
- N variables, N^2 values for each variable representing the positions on the chessboard
- e.g., $Q1 = 1, Q2 = 15, Q3 = 21, Q4 = 32, Q5 = 34, Q6 = 44, Q7 = 54, Q8 = 59$



- This representation has N^{2N} states
 - For 8-Queens: $64^8 = 281,474,976,710,656$
- Is there a better way to represent the N-queens problem?
 - We know we cannot place two queens in a single row, we can exploit this fact

Formulation 2

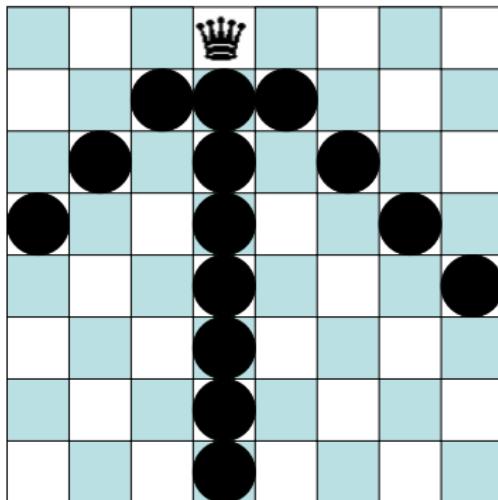
- N variables Q_i , value of Q_i is the column the Queen in row i is placed; possible values $\{1, \dots, N\}$.
- e.g., $Q1 = 1, Q2 = 7, Q3 = 5, Q4 = 8, Q5 = 2, Q6 = 4, Q7 = 6, Q8 = 3$



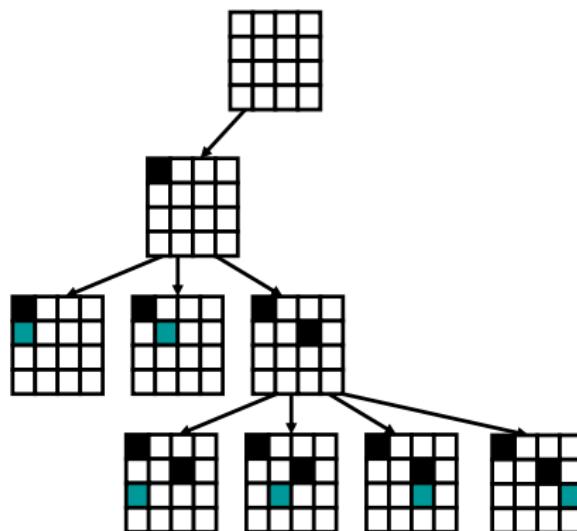
- This representation has N^N states
 - For 8-Queens: $8^8 = 16,777,216$
- The choice of a representation can make the problem solvable or unsolvable!

Constraints

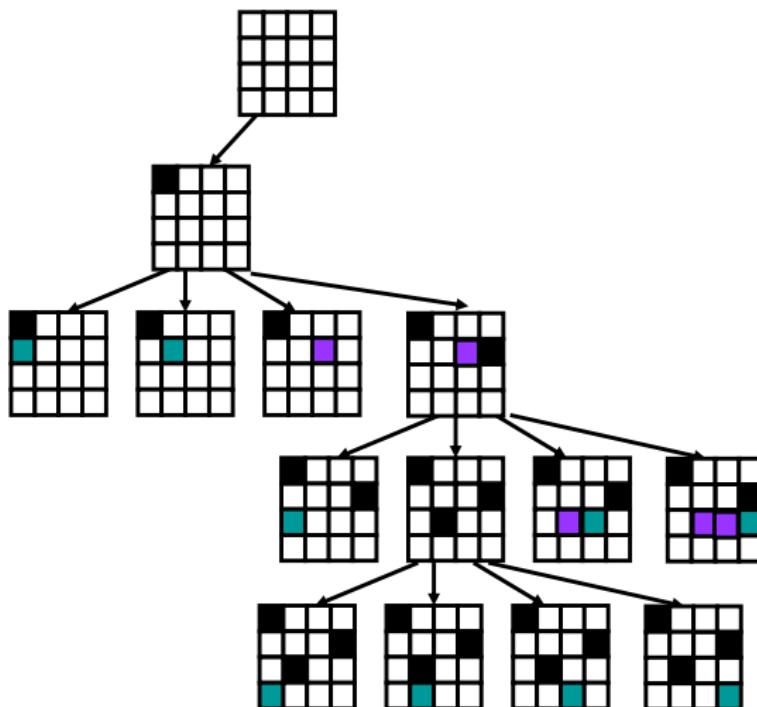
- Can't put two Queens in same column: $Q_i \neq Q_j$ for all $i \neq j$
- Diagonal constraints: $\text{abs}(Q_i - Q_j) \neq \text{abs}(i - j)$



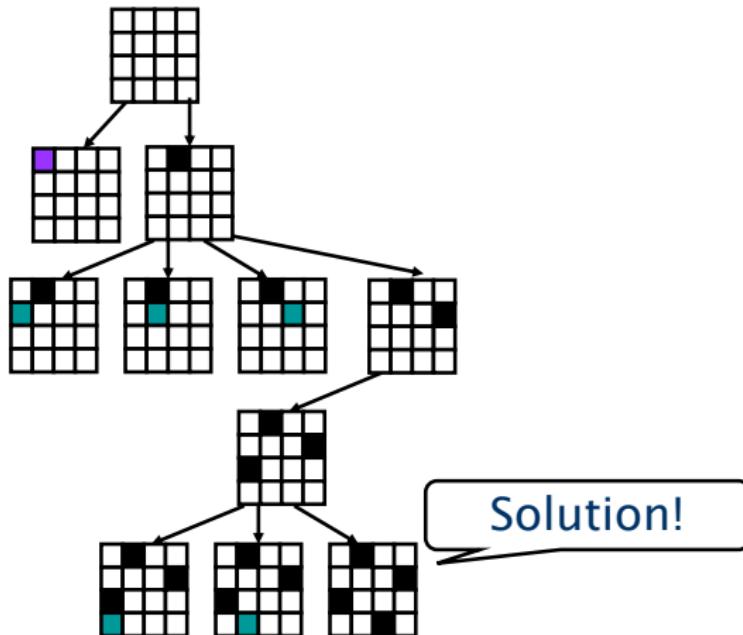
Solving 4-queens



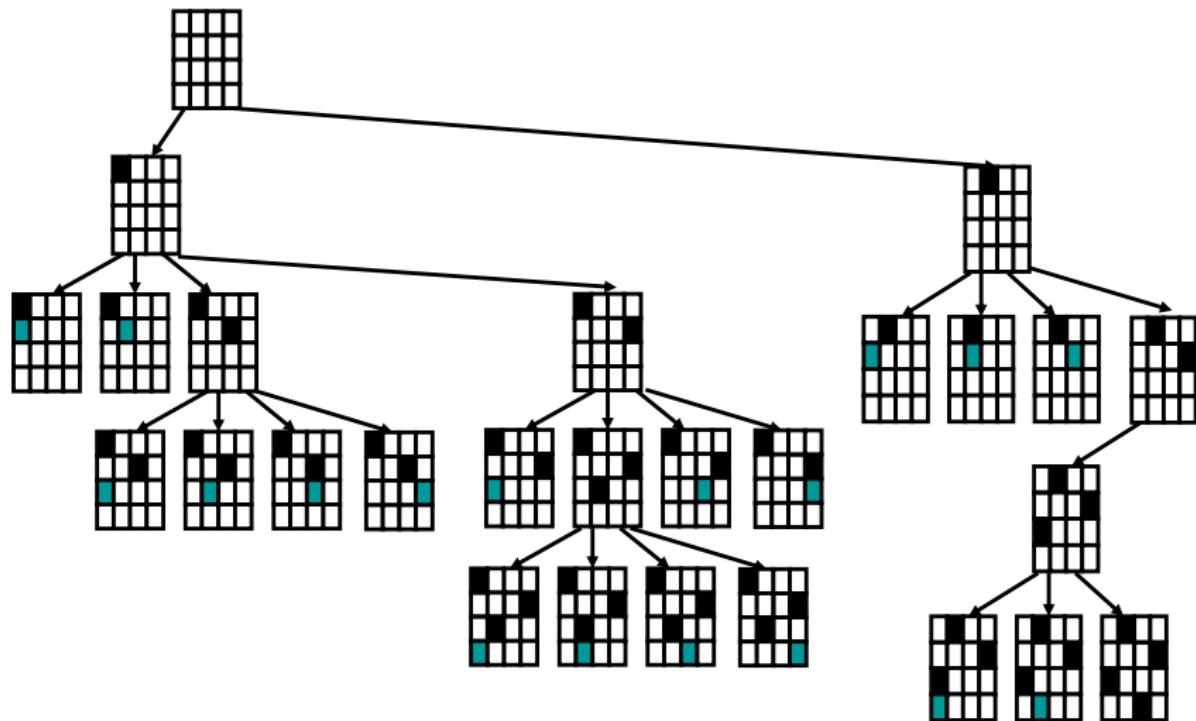
Solving 4-queens



Solving 4-queens



Example: 4-queens backtracking search space



Formalization of a CSP

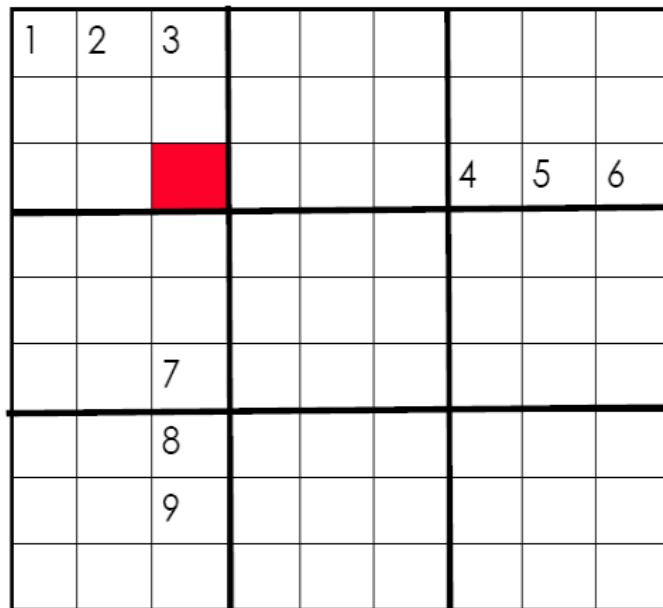
- A CSP consists of
 - A set of variables V_1, \dots, V_n
 - For each variable a (finite) domain of possible values $\text{Dom}[V_i]$.
 - A set of constraints C_1, \dots, C_m .
- A solution to a CSP is an assignment of a value to all of the variables such that every constraint is satisfied.
- A CSP is unsatisfiable if no solution exists.

A generic backtracking algorithm

- We pick a variable*,
- pick a value for it*,
- test the constraints that we can,
- if a constraint is unsatisfied we backtrack,
- otherwise we set another variable.
- When all the variables are set, we're done.

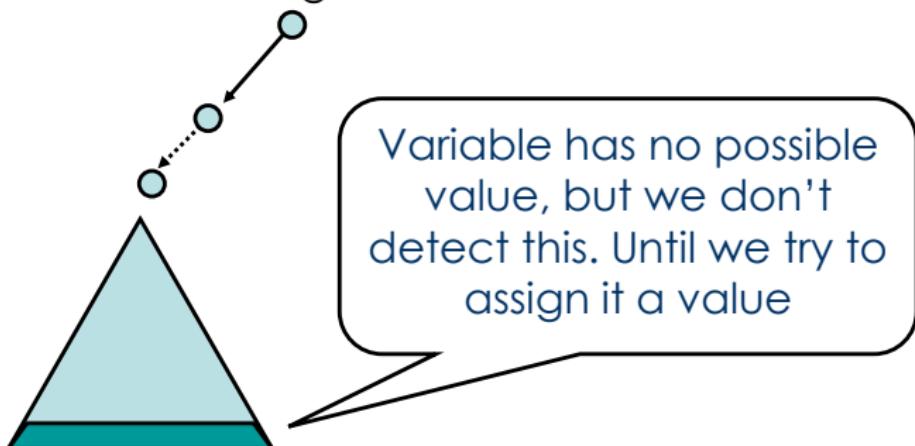
Problems with Plain Backtracking

Sudoku: The 3,3 cell has no possible value.



Problems with Plain Backtracking

- In the backtracking search we won't detect that the (3,3) cell has no possible value until all variables of the row/column (involving row or column 3) or the sub-square constraint (first sub-square) are assigned.
So we have the following situation:



- Leads to the idea of **constraint propagation**

Constraint Propagation (Inference)

- Constraint propagation refers to the technique of “**looking ahead**” at the yet unassigned variables in the search .
- Try to detect obvious failures: “**Obvious**” means things we can test/detect efficiently.
- Even if we don’t detect an obvious failure we might be able to eliminate some possible part of the future search.

Constraint Propagation

- Propagation has to be applied during the search; potentially at every node of the search tree.
- It can also be applied before search even begins!
- Propagation itself is an **inference step** that needs some resources (in particular time)
 - If propagation is slow, this can slow the search down to the point where using propagation makes finding a solution take longer!
 - There is always a **tradeoff** between searching fewer nodes in the search, and having a higher nodes/second processing rate.
- We will look at two main types of propagation:
 1. Forward Checking
 2. Generalized Arc Consistency

Forward Checking

- **Forward checking** is an extension of backtracking search that employs a “modest” amount of propagation (look ahead).
- When a variable is instantiated we check all constraints that have **only one uninstantiated variable** remaining.
- For that uninstantiated variable, we check all of its values, pruning those values that violate the constraint.

Forward Checking Algorithm

For a single constraint C:

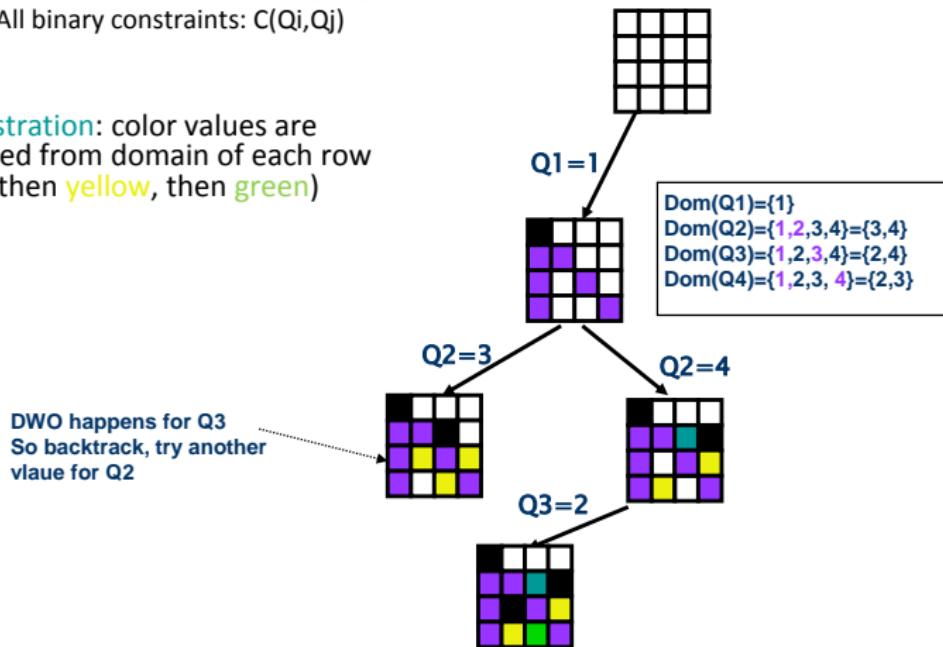
```
FCCheck(C, x)
  // C is a constraint with all its variables already
  // assigned, except for variable x.
  for d := each member of CurDom[x]
    IF making x = d together with previous assignments
      to variables in scope C falsifies C
      THEN remove d from CurDom[x]
  IF CurDom[x] = {} then return DWO (Domain Wipe Out)
  ELSE return ok
```

Forward Checking Algorithm

```
FC(Level) /*Forward Checking Algorithm */
    If all variables are assigned
        PRINT Value of each Variable
        RETURN or EXIT (RETURN for more solutions)
            (EXIT for only one solution)
    V := PickAnUnassignedVariable()
    Assigned[V] := TRUE
    for d := each member of CurDom(V)
        Value[V] := d
        DWOoccurred:= False
        for each constraint C over V such that
            a) C has only one unassigned variable X in its scope
            if(FCCheck(C,X) == DWO) /* X domain becomes empty*/
                DWOoccurred:= True
                break /* stop checking constraints */
            if(not DWOoccurred) /*all constraints were ok*/
                FC(Level+1)
                RestoreAllValuesPrunedByFCCheck()
    Assigned[V] := FALSE //undo since we have tried all of V's values
    return;
```

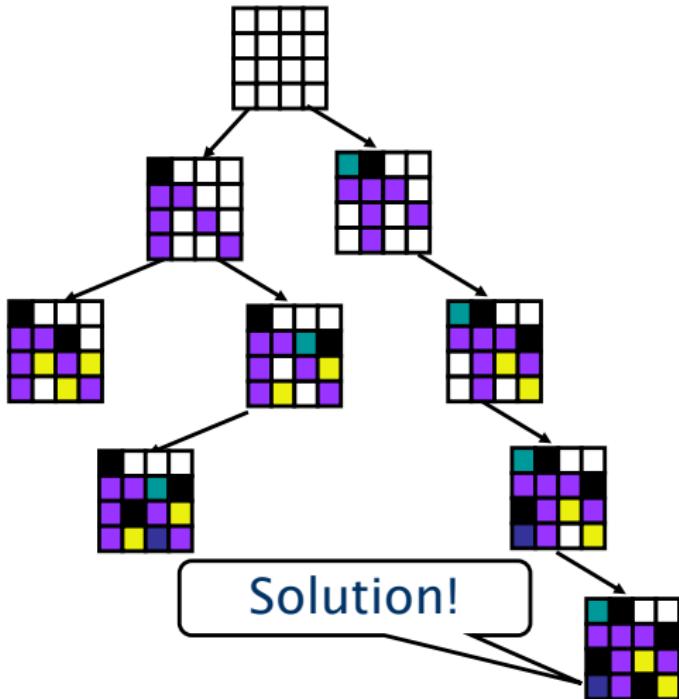
FC 4-queens

- 4X4 Queens
 - Q1,Q2,Q3,Q4 with domain {1..4}
 - All binary constraints: $C(Q_i, Q_j)$
- FC illustration: color values are removed from domain of each row (blue, then yellow, then green)

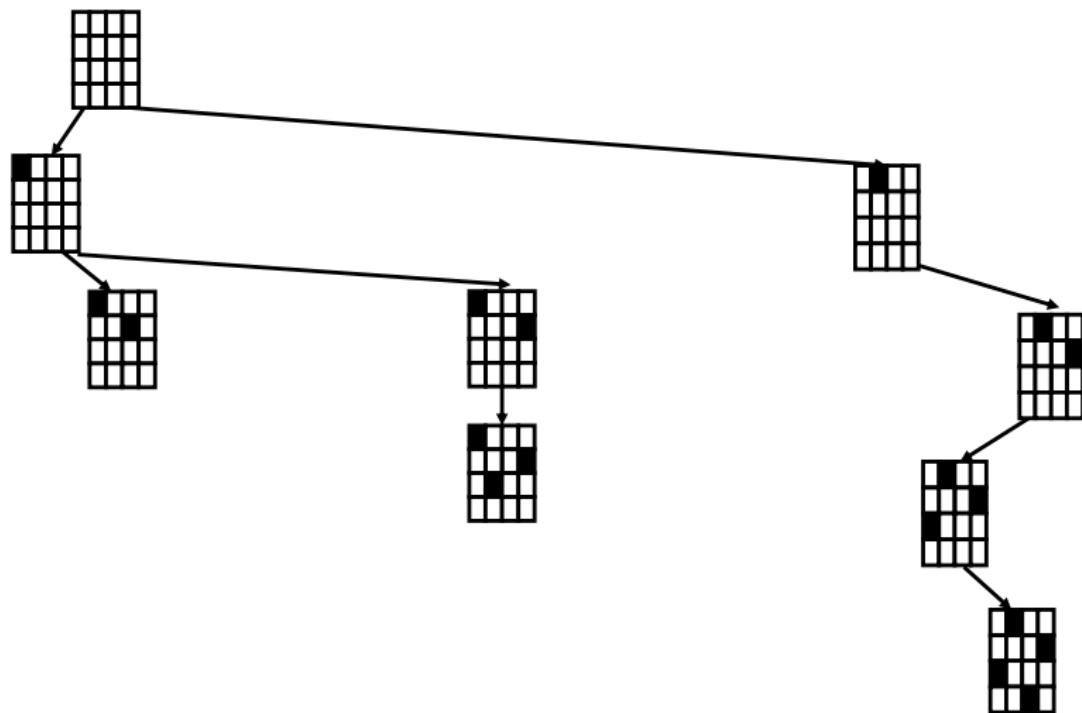


FC 4-queens

- 4X4 Queens
continue...



FC 4-queens search space



FC: Restoring Values

- After we backtrack from the current assignment (in the for loop) we must restore the values that were pruned as a result of that assignment.
- Some bookkeeping needs to be done, as we must remember which values were pruned by which assignment (FCCheck is called at every recursive invocation of FC).

FC: Minimum Remaining Values Heuristics (MRV)

FC also gives us for free a very powerful heuristics regarding which variables to try next:

- Always branch on a variable with the smallest remaining values (smallest CurDom).
- If a variable has only one value left, that value is forced, so we should propagate its consequences immediately.
- This heuristic tends to produce skinny trees at the top.
- This means that more variables can be instantiated with fewer nodes searched, and thus more constraint propagation/DWO failures occur when the tree starts to branch out
- We can find an inconsistency much faster

MRV Heuristic: Human Analogy

- What variables would you try first?

8	1	5	6					4
6				7	5			8
				9				
9				4	1	7		
	4						2	
		6	2	3				8
				5				
	5		9	1				6
1					7	8	9	5

Domain of each variable:
 $\{1, \dots, 9\}$

(1, 5) impossible values:
Row: $\{1, 4, 5, 6, 8\}$
Column: $\{1, 3, 4, 5, 7, 9\}$
Subsquare: $\{5, 6, 7, 9\}$
→ Domain = $\{2\}$

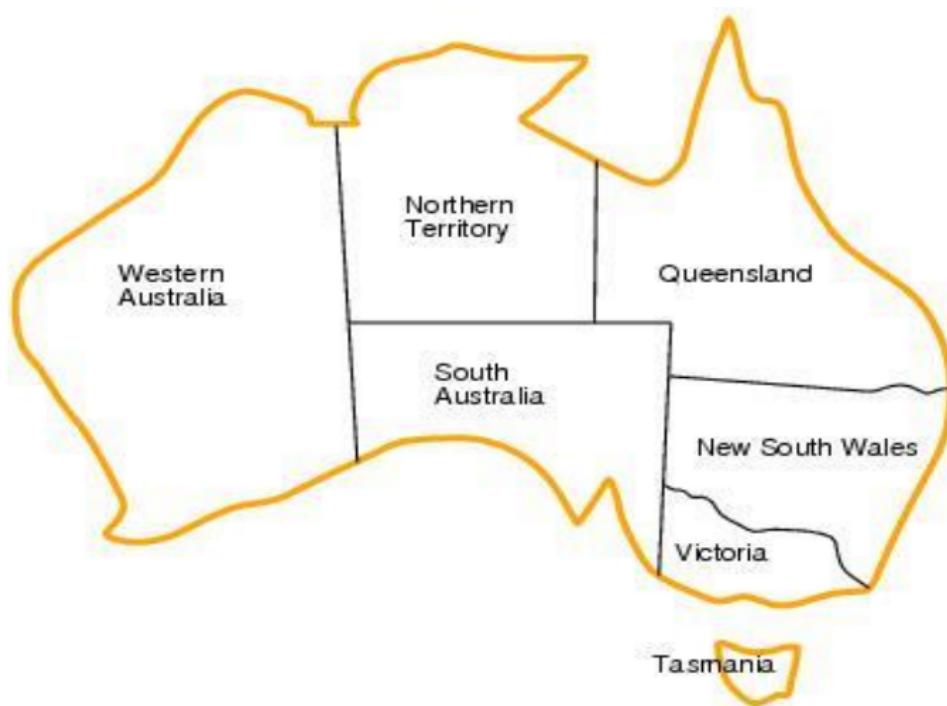
(9, 5) impossible values:
Row: $\{1, 5, 7, 8, 9\}$
Column: $\{1, 3, 4, 5, 7, 9\}$
Subsquare: $\{1, 5, 7, 9\}$
→ Domain = $\{2, 6\}$

After assigning value 2 to
cell (1,5): Domain = {6}

Most restricted variables! = MRV

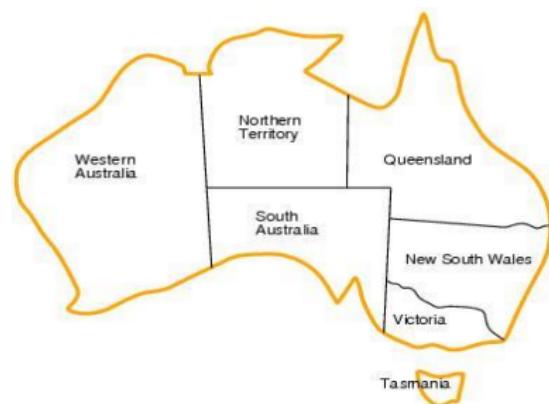
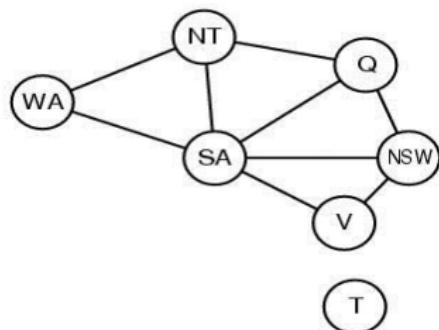
Example – Map coloring

Color the following map using red, green, and blue such that adjacent regions have different colors



Formulation

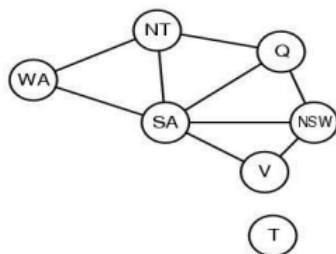
- Variables: WA, NT, Q, NSW, V, SA, T
- Domains: $D_i = \{red, green, blue\}$
- Constraints: adjacent regions must have different colors, e.g., $WA \neq NT$



Example – Map coloring

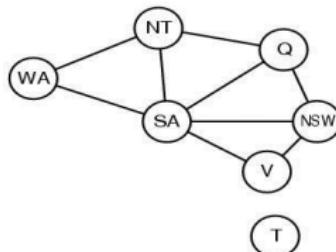
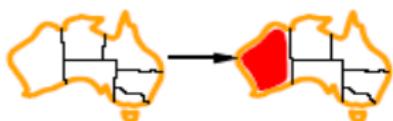
Forward checking idea:

- keep track of remaining legal values for unassigned variables.
- Terminate search when any variable has no legal values.



Example – Map coloring

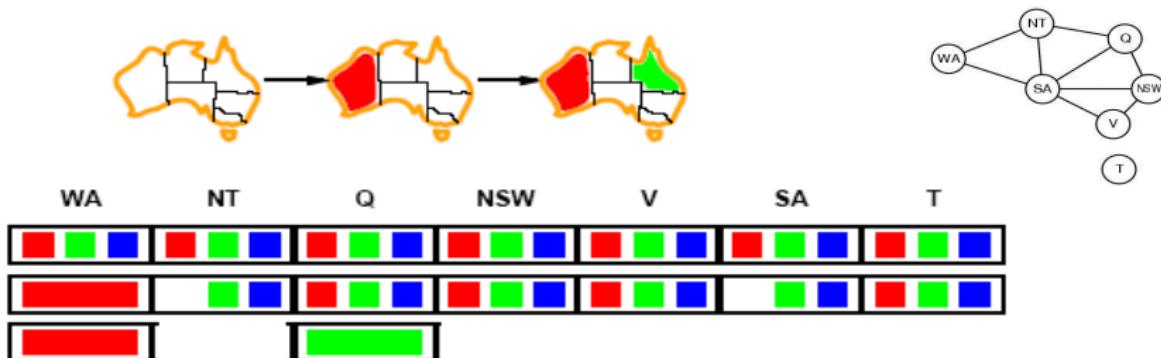
- Assign $\{WA=\text{red}\}$
- Effects on other variables connected by constraints to WA
 - *NT can no longer be red*
 - *SA can no longer be red*



WA NT Q NSW V SA T

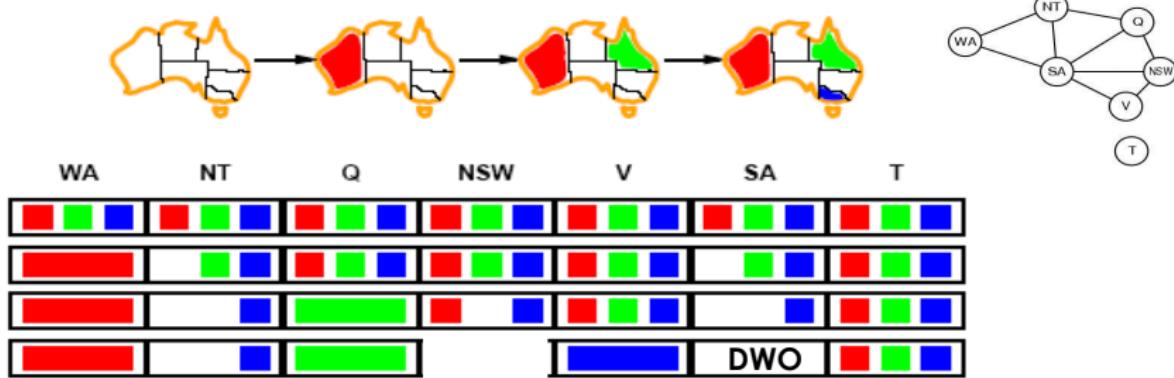
Example – Map coloring

- Assign $\{Q=\text{green}\}$ (Note: Not using MRV)
- Effects on other variables connected by constraints with Q
 - NT can no longer be green
 - NSW can no longer be green
 - SA can no longer be green
- MRV heuristic would automatically select NT or SA next



Example – Map coloring

- Assign $\{V=\text{blue}\}$ (*not using MRV*)
- Effects on other variables connected by constraints with V
 - NSW can no longer be blue
 - SA is empty (DWO – Domain Wipe Out!)
- FC has detected that partial assignment is *inconsistent* with the constraints and backtracking can occur.



Exercise

Try the previous map coloring example with MRV

Empirically

- FC often is about 100 times faster than BT
- FC with MRV (minimal remaining values) often 10000 times faster.
- But on some problems the speed up can be much greater
 - Converts problems that are not solvable to problems that are solvable
- Still FC is not that powerful.
- Other more powerful forms of constraint propagation are used in practice.

Generalized Arc Consistency (GAC)

- $C(X, Y)$ is consistent iff for every value of X there is some value of Y that satisfies C
- $C(V_1, V_2, V_3, \dots, V_n)$ is GAC wrt V_i iff for every value of V_i , there are values of $V_1, \dots, V_{i-1}, V_{i+1}, \dots, V_n$ that satisfy C
- A constraint is GAC iff it is GAC wrt every variable in its scope
- A CSP is GAC iff all of its constraints are GAC

- Say we find a value d of variable V_i that is not consistent wrt a constraint
- i.e., there is no assignments to the other variables that satisfy the constraint when $V_i = d$
- d is said to be arc inconsistent
- We can remove d from the domain of V_i as this value cannot lead to a solution
- Much like Forward Checking, but more powerful
 - If d is removed from $CurDom[V_i]$, it must be the case that d is arc inconsistent

GAC: an example

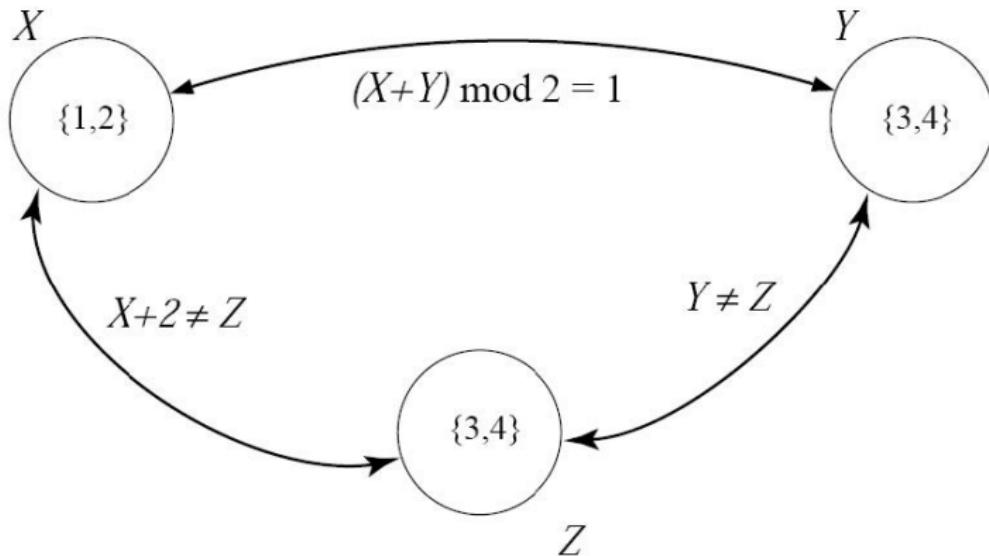
$C(X,Y): X > Y, \text{Dom}(X)=\{1,5,11\}, \text{Dom}(Y)=\{3,8,15\}$

- For $X=1$ there is no value of Y s.t. $1 > Y$, so remove 1 from domain X
- For $Y=15$ there is no value of X s.t. $X > 15$, so remove 15 from domain Y
- We obtain $\text{Dom}(X)=\{5,11\}$ and $\text{Dom}(Y)=\{3,8\}$.

- Removing a value from a domain may trigger further inconsistency, so we have to repeat the procedure until everything is consistent.
- We put constraints on a queue and add new constraints to the queue as we need to check for arc consistency
- If we apply arc consistency propagation during search the search tree will typically be much reduced in size.

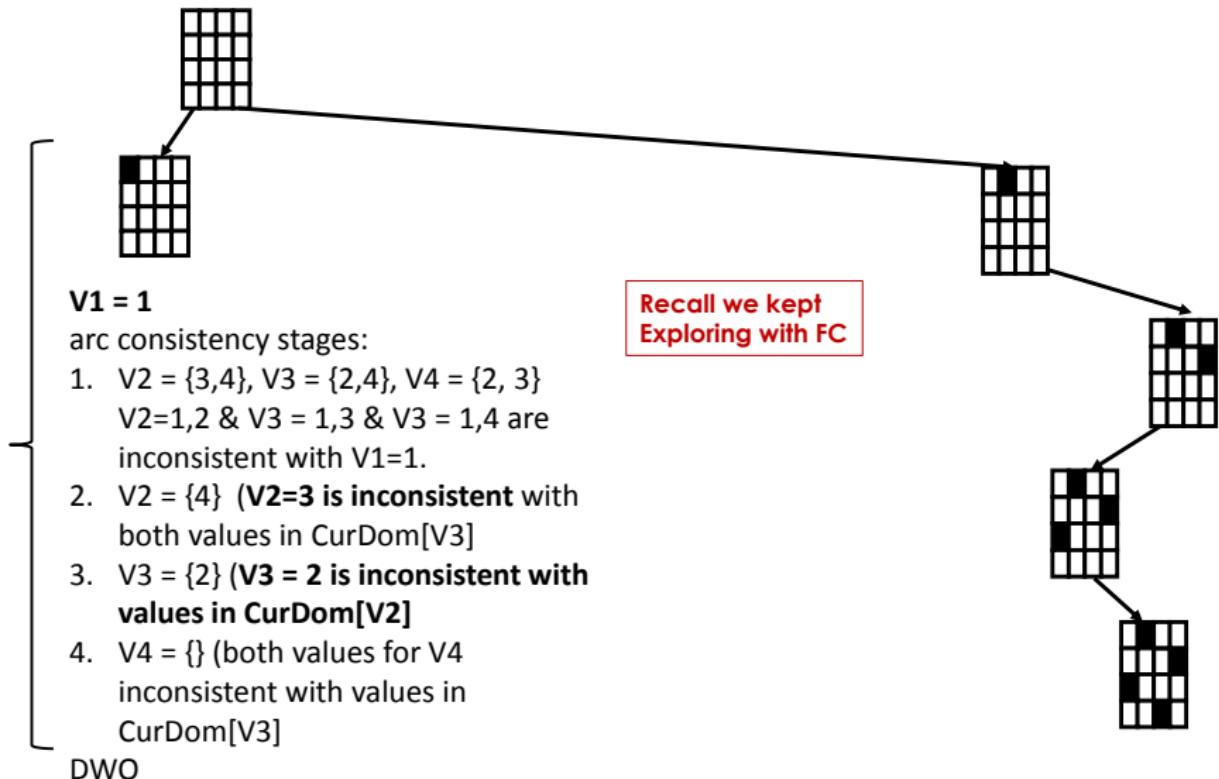
A simple example

Is this CSP arc consistent? If not, what values should be removed?

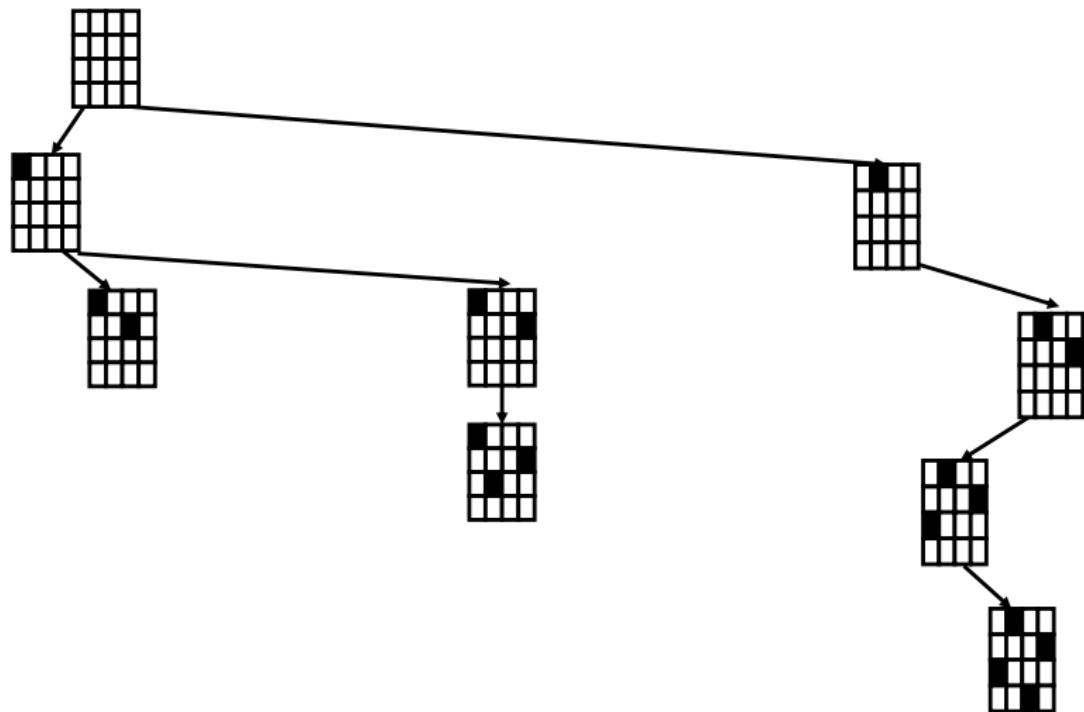


What about $\text{Dom}[Y]=\{2,4\}$?

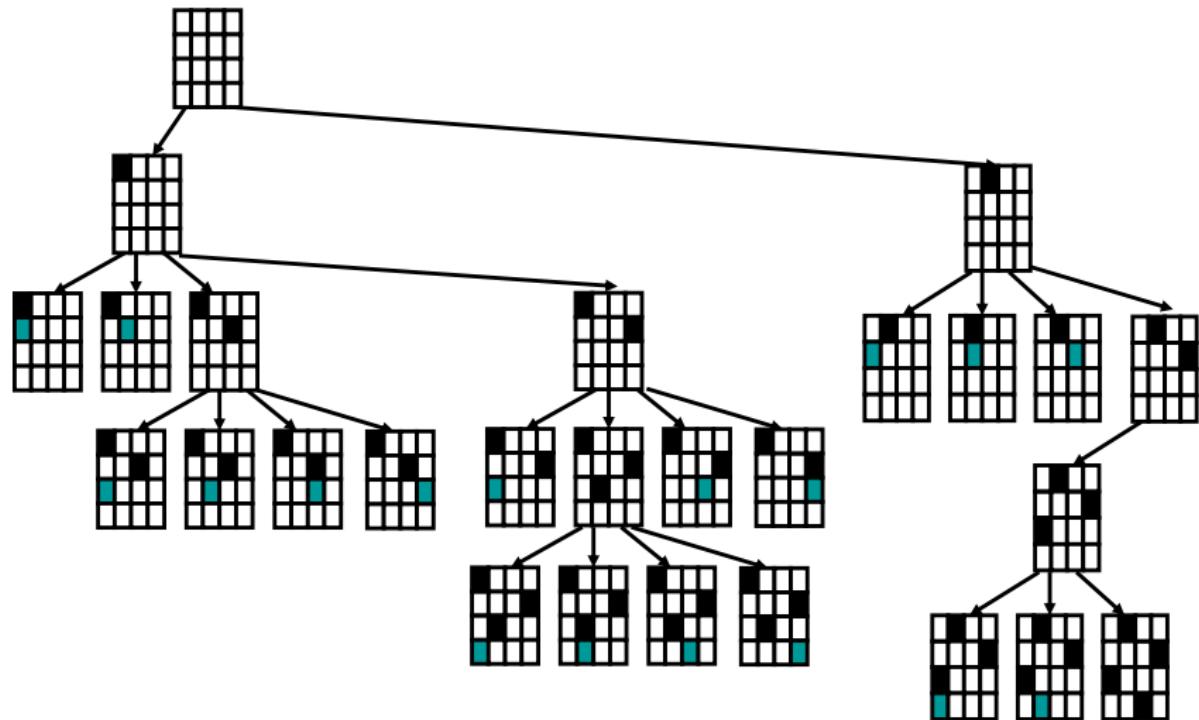
Example: 4-Queens GAC search Space



CONTRAST TO: 4-Queens FC search Space

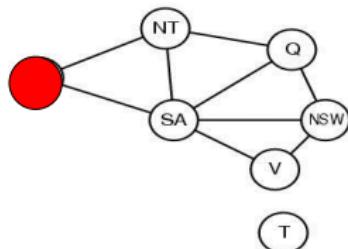


CONTRAST TO: 4-Queens Backtracking Search Space



Example – Map coloring

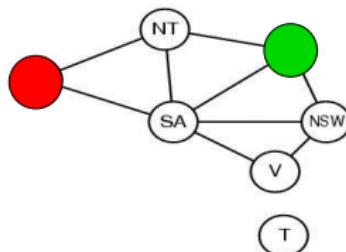
- Assign $\{WA=\text{red}\}$
- Effects on other variables connected by constraints to WA
 - $NT \text{ can no longer be red} = \{G, B\}$
 - $SA \text{ can no longer be red} = \{G, B\}$
- *All other values are arc-consistent*



Example – Map coloring

- Assign $\{Q=\text{green}\}$
- Effects on other variables connected by constraints with Q
 - $NT \text{ can no longer be green} = \{B\}$
 - $NSW \text{ can no longer be green} = \{R, B\}$
 - $SA \text{ can no longer be green} = \{B\}$
- DWO there is no value for SA that will be consistent with $NT \neq SA$ and $NT = B$

Note Forward Checking would not have detected this DWO.



GAC algorithm

- We make all constraints GAC at every node of the search space.
- Like forward checking, GAC could also be performed before we even start to search, *i.e.*, before we assign any variables.
- This is accomplished by removing from the domains of the variables all arc inconsistent values.

GAC Algorithm, enforce GAC during search

```
GAC(Level) /*Maintain GAC Algorithm */
    If all variables are assigned
        PRINT Value of each Variable
        RETURN or EXIT (RETURN for more solutions)
                            (EXIT for only one solution)
    V := PickAnUnassignedVariable()
    Assigned[V] := TRUE
    for d := each member of CurDom(V)
        Value[V] := d
        Prune all values of V ≠ d from CurDom[V]
        for each constraint C whose scope contains V
            Put C on GACQueue
            if(GAC_Enforce() != DWO)
                GAC(Level+1) /*all constraints were ok*/
                RestoreAllValuesPrunedFromCurDoms()
    Assigned[V] := FALSE
    return;
```

Enforce GAC (prune all GAC inconsistent values)

```
GAC_Enforce()
    // GAC-Queue contains all constraints one of whose variables has
    // had its domain reduced. At the root of the search tree
    // first we run GAC_Enforce with all constraints on GAC-Queue
    while GACQueue not empty
        C = GACQueue.extract()
        for V := each member of scope(C)
            for d := CurDom[V]
                Find an assignment A for all other
                variables in scope(C) such that
                C(A  $\cup$  V=d) = True
                if A not found
                    CurDom[V] = CurDom[V] - d
                    if CurDom[V] =  $\emptyset$ 
                        empty GACQueue
                        return DWO //return immediately
                    else
                        push all constraints C' such that
                        V  $\in$  scope(C') and C'  $\notin$  GACQueue
                        on to GACQueue
    return TRUE //while loop exited without DWO
```

Enforce GAC

- A support for $V=d$ in constraint C is an assignment A to all of the other variables in $\text{scope}(C)$ s.t. $A \cup \{V=d\}$ satisfies C .
- Smarter implementations keep track of “supports” to avoid having to search through all possible assignments to the other variables for a satisfying assignment.
- Rather than search for a satisfying assignment to C containing $V=d$, we check to see if the current support is still valid: i.e., all values it assigns still lie in the variable’s current domains
- Also we take advantage that a support for $V=d$, e.g. $\{V=d, X=a, Y=b, Z=c\}$ is also a support for $X=a$, $Y=b$, and $Z=c$

Enforce GAC

- However, finding a support for $V=d$ in constraint C still in the worst case requires $O(2^k)$ work, where k is the arity of C .
- Another key development in practice is that for some constraints this computation can be done in polynomial time.
- e.g., $\text{all-diff}(V_1, \dots, V_n)$ we can check if $V_i = d$ has a support in the current domains of the other variables in polynomial time using ideas from graph theory.

GAC enforce example

8	1	5	6					4
6				7	5			8
				9				
9				4	1	7		
	4							2
	6	2	3					8
			5					
	5		9	1				6
1					7	8	9	5

 = All-diff

$$C_{SS2} = \text{All-diff}(V_{1,4}, V_{1,5}, V_{1,6}, V_{2,4}, V_{2,5}, V_{2,6}, V_{3,4}, V_{3,5}, V_{3,6})$$

$$C_{R1} = \text{All-diff}(V_{1,1}, V_{1,2}, V_{1,3}, V_{1,4}, V_{1,5}, V_{1,6}, V_{1,7}, V_{1,8}, V_{1,9})$$

$$C_{CS} = \text{All-diff}(V_{1,5}, V_{2,5}, V_{3,5}, V_{4,5}, V_{5,5}, V_{6,5}, V_{7,5}, V_{8,5}, V_{9,5})$$

By going back and forth between constraints we get more values pruned.

Consider the case where we perform GAC inference *before* we do any search, limiting ourselves to consideration of only C_{SS2} , C_{R1} and C_{CS} . We place these constraints on the GAC queue and call GAC enforce.

$GAC(C_{SS2}) \rightarrow$

CurDom of $V_{1,5}, V_{1,6}, V_{2,4}, V_{3,4}, V_{3,6} = \{1, 2, 3, 4, 8\}$

$GAC(C_{R1}) \rightarrow$

CurDom of $V_{1,7}, V_{1,8} = \{2, 3, 7, 9\}$

CurDom of $V_{1,5}, V_{1,6} = \{2, 3\}$

$GAC(C_{CS}) \rightarrow$

CurDom of $V_{1,5} = \{2\}$

CurDom of $V_{5,5}, V_{9,5} = \{2, 6, 8\}$

$GAC(C_{SS2}) \rightarrow$

CurDom of $V_{2,4}, V_{3,4}, V_{3,6} = \{1, 4, 8\}$

$GAC(C_{R1}) \rightarrow$

CurDom of $V_{1,6} = \{3\}$

CurDom of $V_{1,7}, V_{1,8} = \{7, 9\}$

Exercise: Forward checking

Consider the CSP with

- Variables = $\{V_1, \dots, V_5\}$.
- $Dom[V_1] = \{1, 2, 3\}$, $Dom[V_2] = \{1, 2, 3\}$,
 $Dom[V_3] = \{1, 2, 3\}$, $Dom[V_4] = \{1, 2, 3, 4, 5\}$,
 $Dom[V_5] = \{1, 2, 3, 4\}$.
- Constraints: $V_i \neq V_j$, $i \neq j$.

Apply forward checking search to find the first solution.

- Instantiate the variables in the order V1, V2, V3, V4, V5.
- Assign the values in the current domain in increasing order.
- At each node indicate the variable being instantiated, and for each variable its updated CurDom.

GAC exercise

- $\text{Dom}[X] = \text{Dom}[Y] = \text{Dom}[Z] = \{1, 2, 3, 4\}$,
 $\text{Dom}[W] = \{1, 2, 3, 4, 5\}$
- $C_1 : X = Y + Z, C_2 : W > X, C_3 : W = X + Y + Z$
- Enforce GAC
- Use GAC algorithm to find the first solution

GAC example

■ C1(V1,V2,V3)

V1	V2	V3
A	B	C
B	A	C
A	A	B

■ C2(V1,V3,V4,V5)

V1	V3	V4	V5
A	A	A	A
A	B	C	B
B	C	B	B
C	A	B	C
C	B	A	B

■ C3(V2,V3,V5)

V2	V3	V5
A	A	A
A	B	C
B	C	B
C	A	B
C	B	A

■ Dom[V1]...Dom[V5] = {a, b, c}