

Search

- Problem solving by search
- Uninformed search
- Heuristic search

Readings: Chapter 3

*Slides based on those of Sheila McIlraith

Why search

- Success in game playing programs based on search
- Many other AI problems can be successfully solved by search

The formalism

To formulate a problem as a search problem we need the following components:

- ① Formulate a **state space** over which to search. The state space necessarily involves abstracting the real problem.
- ② Formulate **actions** that allow one to move between different states. The actions are abstractions of actions you could actually perform.
- ③ Identify the **initial state** that best represents your current state
- ④ Identify the **goal** or **desired condition** one wants to achieve.
- ⑤ Formulate various **heuristics** to help guide the search process.

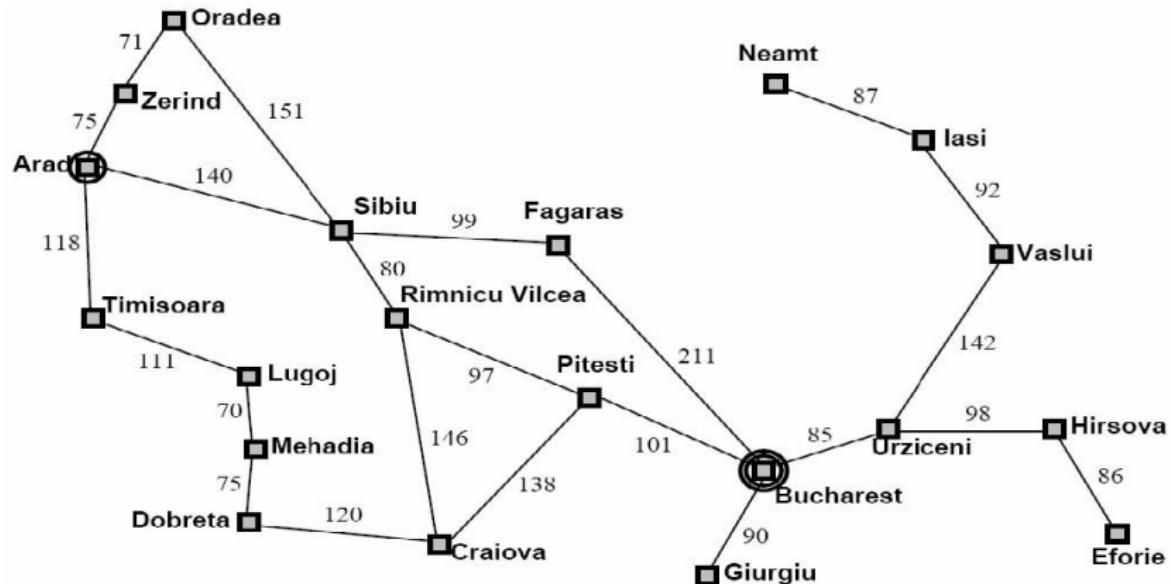
The formalism

Once the problem has been formulated as a state space search, various algorithms can be utilized to solve the problem.

- A solution to the problem will be a sequence of actions/moves that can transform your current state into state where your desired condition holds.

Example 1: Romania Travel

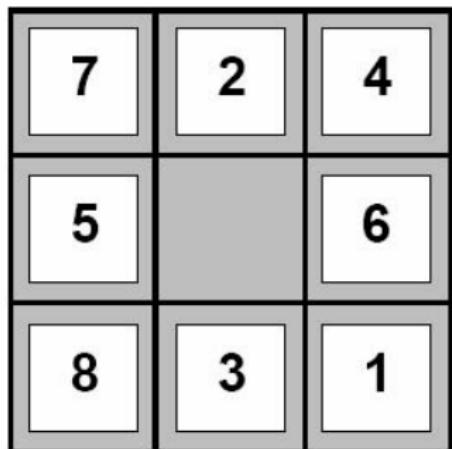
Currently in Arad, need to get to Bucharest



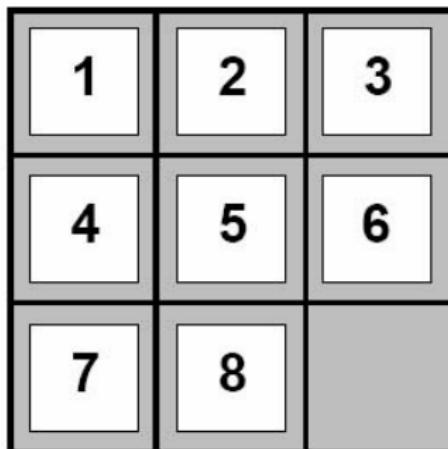
Example 1

- States: the various cities you could be located in.
- Actions: drive between neighboring cities.
- Initial state: in Arad
- Goal: in Bucharest
- Solution: the route, the sequence of cities to travel through to get to Bucharest.

Example 2: The 8-Puzzle



Start State



Goal State

Rule: Can slide a tile into the blank spot. (Equivalently, can think of it as moving the blank around).

Example 2

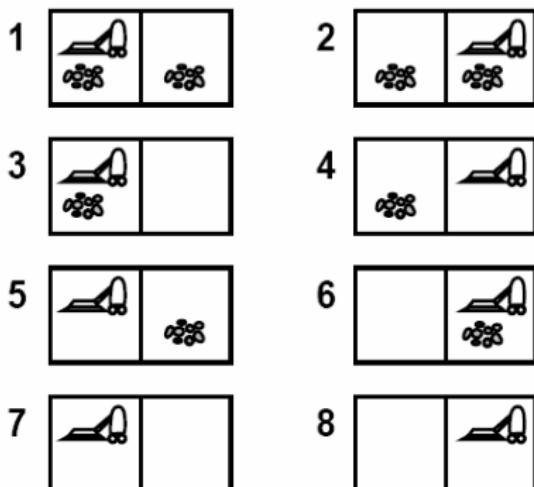
- States: The different configurations of the tiles. How many different states?
- Actions: Moving the blank up, down, left, right. Can every action be performed in every state?
- Initial state: as shown on previous slide.
- Goal: be in the goal state shown on the previous slide.
- Solution: a sequence of moves that transform the initial state to a goal state.

Example 3

- In the previous two examples, a state in the search space corresponded to a unique state of the world
- However, a state could also correspond to the agent's mental conception of how the world is configured: the agent's knowledge state.

Example 3: Vacuum world

- We have a vacuum cleaner and two rooms.
- Each room may or may not be dirty.
- The vacuum cleaner can move **left** or **right** (*the action has no effect if there is no room to the right/left*).
- The vacuum cleaner can **suck**; this cleans the room (*even if the room was already clean*).

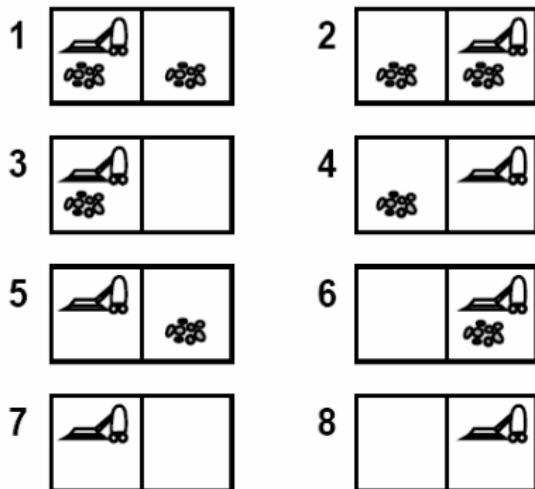


Physical states

Example 3: Vacuum world

Knowledge-level State Space

- The state space can consist of a set of states. The agent knows that it is in one of these states, but doesn't know which.

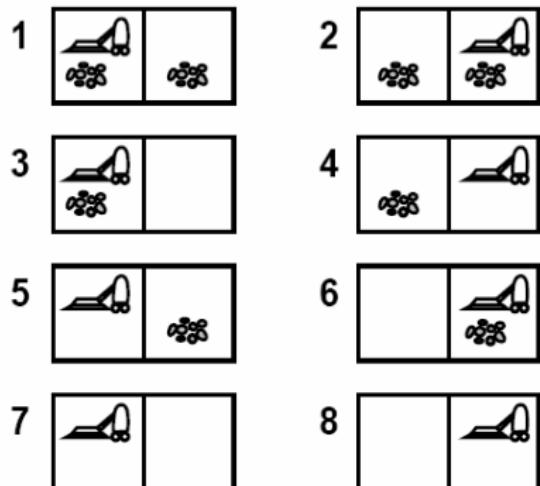


Goal is to have all rooms clean.

Example 3: Vacuum world

Knowledge-level State Space

- **No knowledge of the world.**
States consist of sets of physical states.
- **Start in {1,2,3,4,5,6,7,8},**
agent doesn't have any knowledge of where it is.
- Nevertheless, the actions **<right, suck, left, suck>** achieves the goal.



Goal is to have all rooms clean.

right {2, 4, 6, 8} suck {4, 8} left {3, 7} suck {7}

Inputs:

- a specified **initial state** (a specific world state or a set of world states representing the agent's knowledge, etc.)
- a **successor function** $S(x) = \{\text{set of states that can be reached from state } x \text{ via a single action}\}$.
- a **goal test** a function that can be applied to a state and returns true if the state is satisfies the goal condition.
- A **step cost** function $C(x,a,y)$ which determines the cost of moving from state x to state y using action a . ($C(x,a,y) = \infty$ if a does not yield y from x)

Output:

- a sequence of states leading from the initial state to a state satisfying the goal test.

Obtaining the action sequence

- The set of successors of a state x might arise from different actions, e.g.,
 - $x \rightarrow a \rightarrow y$
 - $x \rightarrow b \rightarrow z$
- Successor function $S(x)$ yields a set of states that can be reached from x via a (any) single action.
 - Rather than just return a set of states, we might annotate these states by the action used to obtain them:
 - $S(x) = \{<y, a>, <z, b>\}$
 y via action a , z via action b .
 - $S(x) = \{<y, a>, <y, b>\}$
 y via action a , also y via alternative action b .

Tree search

- Frontier is the set of states we haven't yet explored/expanded, and want to explore
- Initial call has Frontier = the set of initial state

```
TreeSearch(Frontier, Successors, Goal?)
```

```
If Frontier is empty return failure
```

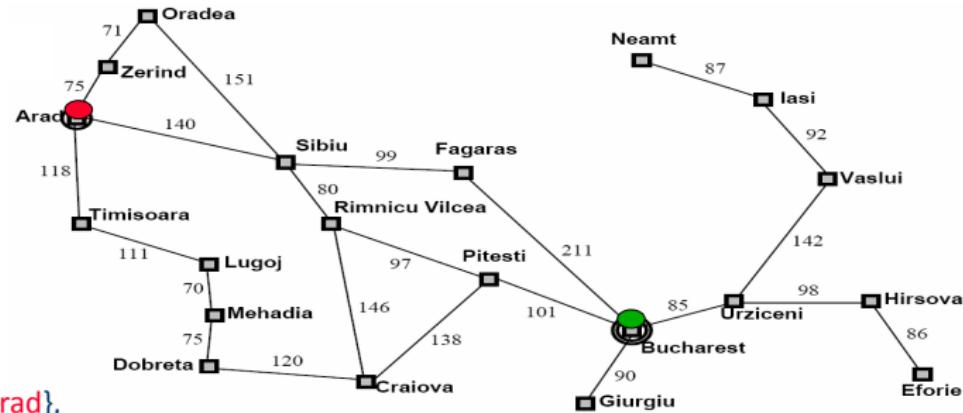
```
Curr = select state from Frontier
```

```
If (Goal?(Curr)) return Curr.
```

```
Frontier' = (Frontier - {Curr}) U Successors(Curr)
```

```
return TreeSearch(Frontier', Successors, Goal?)
```

Example: Romania travel



{Arad},
{Z<A>, T<A>, S<A>},

{Z<A>, T<A>, A<S;A>, O<S;A>, F<S;A>, R<S;A>}

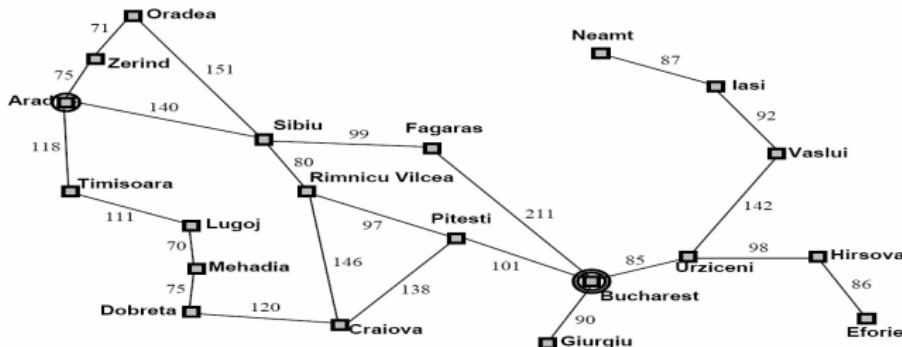
{Z<A>, T<A>, A<S;A>, O<S;A>, R<S;A>, S<F;S;A>, B<F;S;A>}

Solution: Arad -> Sibiu -> Fagaras -> Bucharest

Cost: 140 + 99 + 211 = 450



Another solution



{Arad}

{Z<A>, T<A>, **S<A>**},

{Z<A>, T<A>, A<S,A>, O<S,A>, F<S,A>, **R<S,A>**}

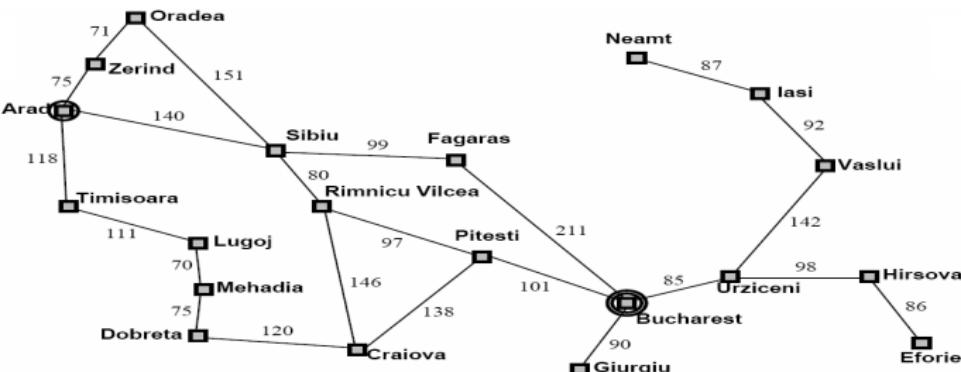
{Z<A>, T<A>, A<S,A>, O<S,A>, F<S,A>, **S<R,S,A>**, **P<R,S,A>**, C<R,S,A>}

{Z<A>, T<A>, A<S,A>, O<S,A>, F<S,A>, S<R,S,A>, C<R,S,A>, R<P,R,S,A>, C<P,R,S,A>, **Bucharest<P,R,S,A>**}

Solution: Arad -> Sibiu -> Rimnicu Vilcea -> Pitesti -> Bucharest

Cost: 140 + 80 + 97 + 101 = 418

The cycle issue



{Arad}

{Z<A>, T<A>, S<A>},

{Z<A>, T<A>, O<S;A>, F<S;A>, A<S;A>, R<S;A>}

{Z<A>, T<A>, O<S;A>, F<S;A>, R<S;A>, Z<A;S;A>, T<A;S;A>, S<A,S,A>}

.....

Frontier is a set of paths not a set of states: **cycles** become an issue.

Selection rule

The example shows that the order states are selected from the frontier has a critical effect on the operation of the search:

- Whether or not a solution is found
- The cost of the solution found
- The time and space required by the search

Critical Properties of Search

- Completeness: will the search always find a solution if a solution exists?
- Optimality: will the search always find the least cost solution? (when actions have costs)
- Time complexity: what is the maximum number of nodes than can be expanded or generated?
- Space complexity: what is the maximum number of nodes that have to be stored in memory?

Uninformed Search Strategies

- These are strategies that adopt a fixed rule for selecting the next state to be expanded.
- The rule does not change irrespective of the search problem being solved.
- These strategies do not take into account any domain specific information about the particular search problem.

Popular uninformed search techniques

- Breadth-First
- Uniform-Cost
- Depth-First
- Depth-Limited
- Iterative-Deepening search

Selecting vs. Sorting

- A simple uniform method we will exploit
 - Order the elements on the frontier.
 - Always select the first element.
- Any selection rule can be achieved by employing an appropriate ordering of the frontier set.

Breadth first

- Place the successors of the current state at the **end** of the frontier.
- Example:
 - let the states be the positive integers $\{0,1,2,\dots\}$
 - let each state n have as successors $n+1$ and $n+2$
 - E.g. $S(1) = \{2, 3\}; S(10) = \{11, 12\}$
 - Start state 0
 - Goal state 5

Example

{0<>}

{1,2}

{2,2,3}

{2,3,3,4}

{3,3,4,3,4}

{3,4,3,4,4,5}

...

Breadth First Properties

- Let b be the maximum number of successors of any state.
- Let d be the number of actions in the shortest solution.

Completeness and optimality

- All shorter paths are expanded before any longer path
- There are finitely many paths of a certain length
- Eventually we must examine all paths of length d , and thus find the shortest solution

Time and space complexities

- Time complexity: $1 + b + b^2 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$
- Space complexity: $b(b^d - 1) = O(b^{d+1})$

Space complexity is a real problem

- E.g., let $b = 10$, and say 1000 nodes can be expanded per second and each node requires 100 bytes of storage:

Depth	Nodes	Time	Memory
1	1	1 millisec.	100 bytes
6	10^6	18 mins.	111 MB
8	10^8	31 hrs.	11 GB

- Run out of space long before we run out of time in most applications.

Depth first

- Place the successors of the current state at the front of the frontier
- Therefore always expands the deepest node in the frontier

Applied to the example of Breadth First Search

{0}

{1,2}

{2,3,2}

{3,4,3,2}

{4,5,4,3,2}

{5,6 4,5,4,3,2}

...

Depth first properties

- Completeness:
 - Infinite state space: No
 - Finite state space with infinite paths: No
 - Finite state space and prune paths with duplicate states ? Yes
- Optimality: No

Time complexity

- $O(b^m)$ where m is the length of the longest path in the state space (Could explore each branch of search tree)
- Very bad if m is much larger than d
- But if there are many solution paths it can be much faster than breadth first (Can by good luck bump into a solution quickly).

Space complexity

- Depth-First Backtrack Points = unexplored siblings of nodes along current path
- Only explore a single path at a time.
- The Frontier only contains the deepest node on the current path along with the backtrack points.
- $O(bm)$, linear space!
- A significant advantage of DFS

Uniform-cost

- Keep Frontier ordered by increasing cost of the path
- Always expand the least cost path
- Identical to breadth first if each action has the same cost.

Completeness and optimality

- Suppose that each transition has costs $\geq \epsilon > 0$
- All cheaper paths are expanded before any more expensive path
- There are finitely many path costs less than the cost of the optimal solution
- Eventually we must examine the optimal solution

Time and space complexities

- Recall the time and space complexity for breadth-first search is $O(b^{d+1})$ where d is the length of the optimal solution
- $O(b^{C^*/\epsilon+1})$ where C^* is the cost of the optimal solution

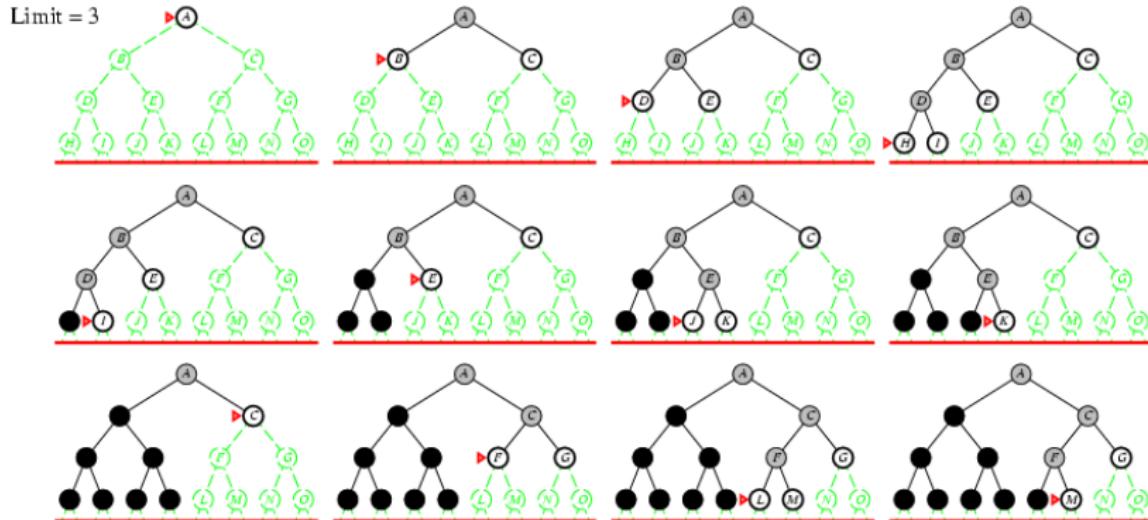
Depth-limited search

- Breadth-first has problems with space complexity
- Depth-first can run down a very long or infinite path

Depth-limited search

- Perform depth first search but only to a pre-specified depth limit L
- Now infinite length paths are not a problem
- But will only find a solution if a solution of length $\leq L$ exists

An example



Depth-limited properties

- Completeness: No
- Optimality: No
- Time complexity: $O(b^L)$
- Space complexity: $O(bL)$

Iterative deepening search

- Solve the problems of depth-first and breadth-first by extending depth limited search
- Starting at depth limit $L = 0$, we iteratively increase the depth limit, performing a depth limited search for each depth limit
- Stop if a solution is found, or if the depth limited search failed without cutting off any nodes because of the depth limit
- If no nodes were cut off, the search examined all paths in the state space and found no solution, hence no solution exists.

An example

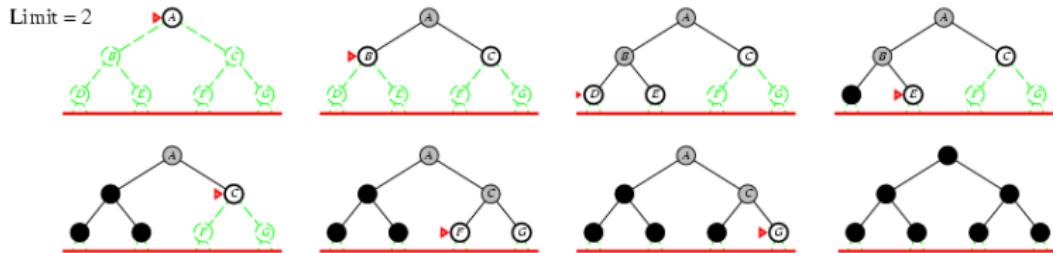
Limit = 0



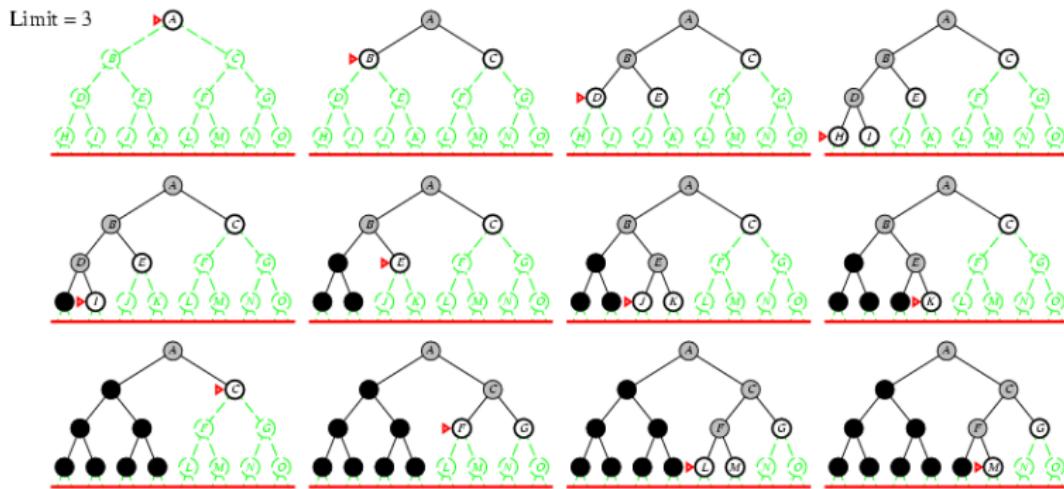
An example



An example



An example



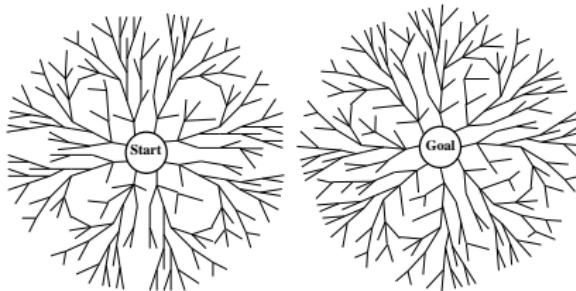
Completeness and optimality

- Completeness: Yes
- Optimality: Yes if costs are uniform
- If costs are not uniform, we can use a cost bound instead
 - Only expand paths of cost less than the cost bound.
 - Keep track of the minimum cost unexpanded path in each depth first iteration, increase the cost bound to this on the next iteration.
 - This can be very expensive. Need as many iterations of the search as there are distinct path costs.

Time and space complexities

- $(d + 1)b^0 + db + (d - 1)b^2 + \dots + b^d = O(b^d)$
- Recall the time complexity for breadth-first:
 $1 + b + b^2 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$
- IDS can be more efficient than breadth first search: nodes at limit are not expanded. BFS must expand all nodes until it expand a goal node
- Space complexity: $O(bd)$

Bidirectional search



- Simultaneously search both forward from the initial state and backward from the goal, and stop when the two searches meet in the middle
- Suppose both directions use breadth-first
- Completeness: Yes
- Optimality: if edges have uniform costs
- Time and space complexity: $O(b^{d/2})$

Summary of uninformed search

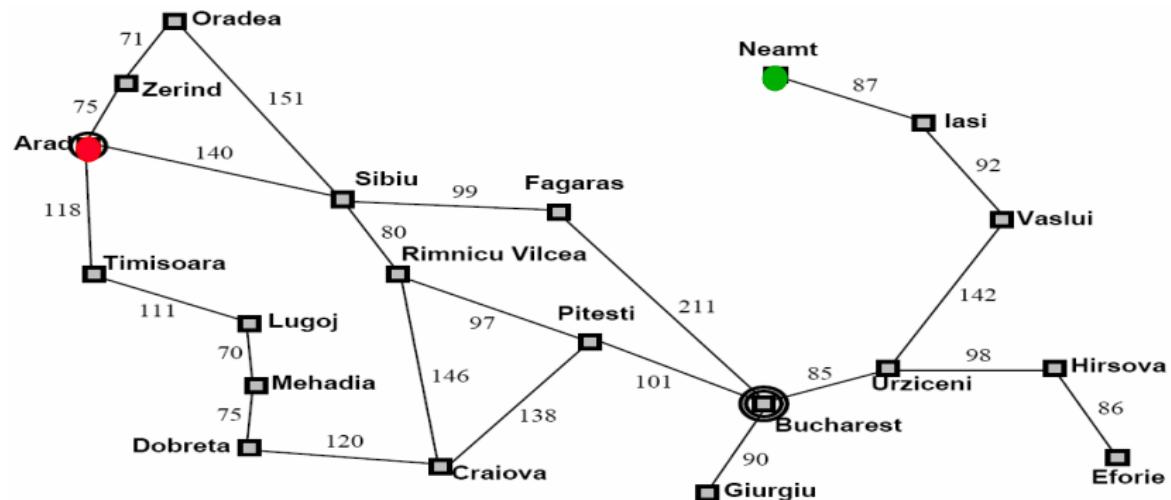
Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

Figure 3.21 Evaluation of tree-search strategies. b is the branching factor; d is the depth of the shallowest solution; m is the maximum depth of the search tree; ℓ is the depth limit. Superscript caveats are as follows: ^a complete if b is finite; ^b complete if step costs $\geq \epsilon$ for positive ϵ ; ^c optimal if step costs are all identical; ^d if both directions use breadth-first search.

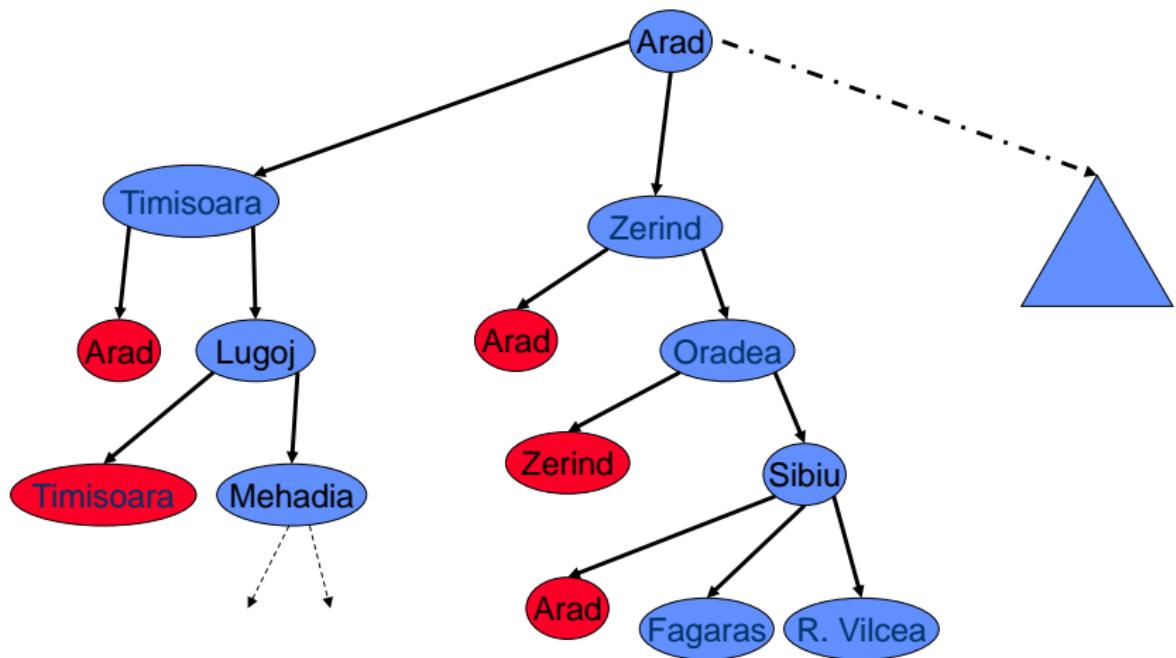
Path checking

- Recall paths are stored on the frontier
- If $\langle n_1, \dots, n_k \rangle$ is a path to node n_k , and we expand n_k to obtain child c , we have $\langle n_1, \dots, n_k, c \rangle$ as the path to c
- Path checking ensures that the state c is not equal to the state reached by any ancestor of c along this path
- That is, paths are checked in isolation!

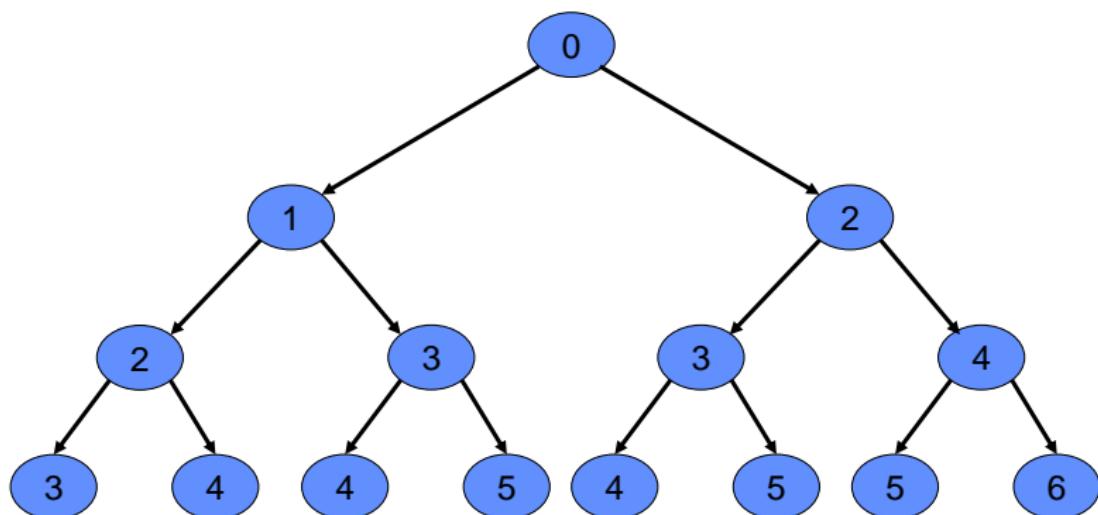
Example: Arad to Neamt



Path checking example



Path checking example

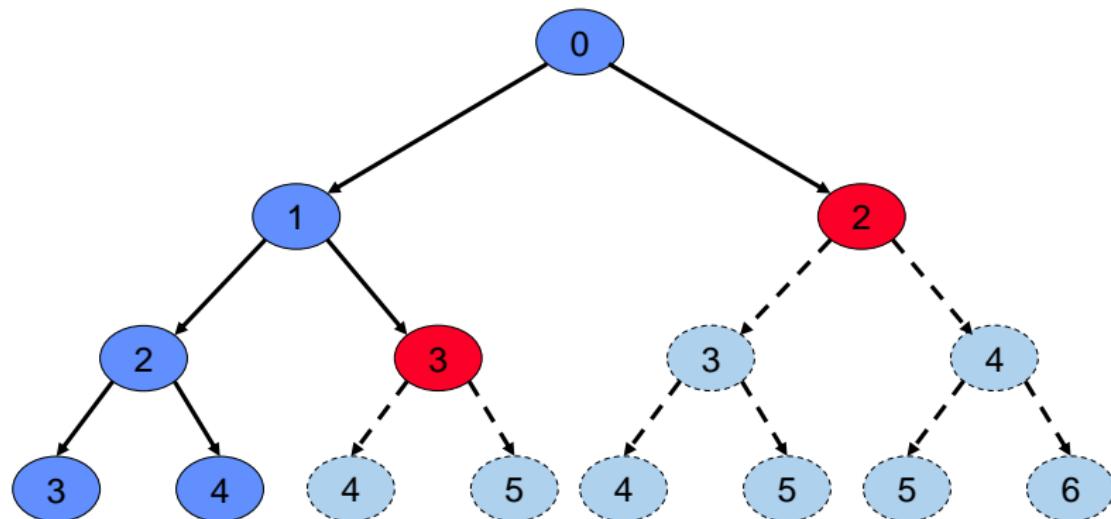


Cycle checking / multiple path checking

- Keep track of all states previously expanded during the search
- When we expand n_k to obtain child c , ensure that c is not equal to any previously expanded state
- Why can't we utilize this technique with depth-first search?
- High space complexity, only useful with breadth first search

Cycle checking example

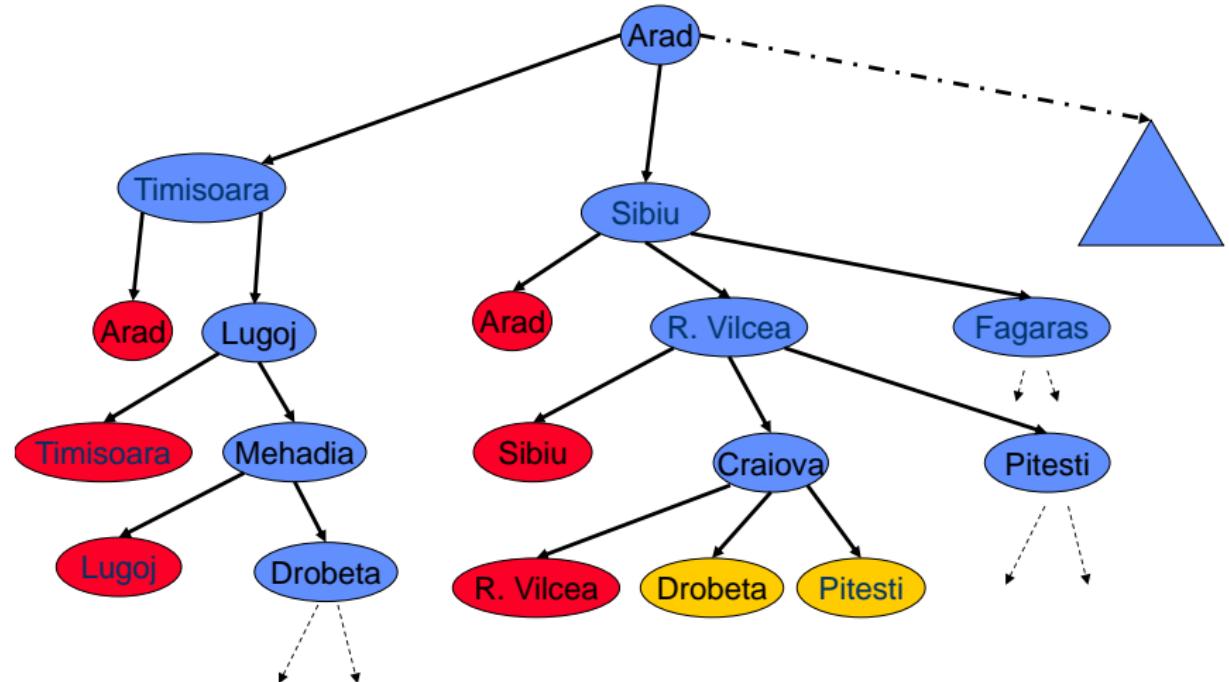
Depth-first search



Issue with optimality

- With uniform-cost search, we still find an optimal solution
- The first time uniform-cost expands a state it has found the minimal cost path to it.
- This means that the nodes rejected by cycle checking can't have better paths.
- We will see later that we don't always have this property when we do heuristic search.

Cycle checking example (BFS)



Summary of uninformed search

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

Figure 3.21 Evaluation of tree-search strategies. b is the branching factor; d is the depth of the shallowest solution; m is the maximum depth of the search tree; ℓ is the depth limit. Superscript caveats are as follows: ^a complete if b is finite; ^b complete if step costs $\geq \epsilon$ for positive ϵ ; ^c optimal if step costs are all identical; ^d if both directions use breadth-first search.

Path / cycle checking

- Path checking: when we expand n to obtain child c , ensures that the state c is not equal to the state reached by any ancestor of c along this path
- Cycle checking: keep track of all states previously expanded during the search; when we expand n to obtain child c , ensure that c is not equal to any previously expanded state
- For uniform-cost search, cycle checking preserves optimality

The missionaries and cannibals problem

- N missionaries and N cannibals are at the left bank of a river
- There is a boat that can hold K people
- Find a way to get everyone to the right bank
- So that at any time, at any place (on either bank, or in the boat), $\#\text{missionaries} \geq \#\text{cannibals}$ or $\#\text{missionaries} = 0$

Formulation of the MC problem

- States (M, C, B) where M – #missionaries, C – #cannibals at the left bank, $B = 1$ indicates the boat is at the left bank
- Actions (m, c) where m – #missionaries, c – #cannibals on the boat
- Precondition: #missionaries and #cannibals satisfy the constraint
- Effects: $(M, C, 1) \xrightarrow{(m,c)} (M - m, C - c, 0)$ and $(M, C, 0) \xrightarrow{(m,c)} (M + m, C + c, 1)$

Exercise

Running breadth-first with cycle-checking for $M = 3$ and $K = 2$