
E10 Decision Tree

18340052

November 25, 2020

Contents

1	Datasets	2
2	Decision Tree	2
2.1	ID3	2
2.2	C4.5 and CART	3
3	Tasks	4
4	Codes and Results	4

1 Datasets

The UCI dataset (<http://archive.ics.uci.edu/ml/index.php>) is the most widely used dataset for machine learning. If you are interested in other datasets in other areas, you can refer to <https://www.zhihu.com/question/63383992/answer/222718972>.

Today's experiment is conducted with the **Adult Data Set** which can be found in <http://archive.ics.uci.edu/ml/datasets/Adult>.

Data Set Characteristics:	Multivariate	Number of Instances:	48842	Area:	Social
Attribute Characteristics:	Categorical, Integer	Number of Attributes:	14	Date Donated	1996-05-01
Associated Tasks:	Classification	Missing Values?	Yes	Number of Web Hits:	1305515

You can also find 3 related files in the current folder, `adult.name` is the description of **Adult Data Set**, `adult.data` is the training set, and `adult.test` is the testing set. There are 14 attributes in this dataset:

>50K, <=50K.

Prediction task is to determine whether a person makes over 50K a year.

2 Decision Tree

2.1 ID3

ID3 (Iterative Dichotomiser 3) was developed in 1986 by Ross Quinlan. The algorithm creates a multiway tree, finding for each node (i.e. in a greedy manner) the categorical feature that will yield the largest information gain for categorical targets. Trees are grown to their maximum size and then a pruning step is usually applied to improve the ability of the tree to generalise to unseen data.

ID3 Algorithm:

1. Begins with the original set S as the root node.
2. Calculate the entropy of every attribute a of the data set S .
3. Partition the set S into subsets using the attribute for which the resulting entropy after splitting is minimized; or, equivalently, information gain is maximum.
4. Make a decision tree node containing that attribute.
5. Recur on subsets using remaining attributes.

Recursion on a subset may stop in one of these cases:

- every element in the subset belongs to the same class; in which case the node is turned into a leaf node and labelled with the class of the examples.
- there are no more attributes to be selected, but the examples still do not belong to the same class. In this case, the node is made a leaf node and labelled with the most common class of the examples in the subset.
- there are no examples in the subset, which happens when no example in the parent set was found to match a specific value of the selected attribute.

ID3 shortcomings:

- ID3 does not guarantee an optimal solution.
- ID3 can overfit the training data.
- ID3 is harder to use on continuous data.

Entropy:

Entropy $H(S)$ is a measure of the amount of uncertainty in the set S .

$$H(S) = \sum_{x \in X} -p(x) \log_2 p(x)$$

where

- S is the current dataset for which entropy is being calculated
- X is the set of classes in S
- $p(x)$ is the proportion of the number of elements in class x to the number of elements in set S .

Information gain:

Information gain $IG(A)$ is the measure of the difference in entropy from before to after the set S is split on an attribute A . In other words, how much uncertainty in S was reduced after splitting set S on attribute A .

$$IG(S, A) = H(S) - \sum_{t \in T} p(t)H(t) = H(S) - H(S | A)$$

where

- $H(S)$ is the entropy of set S
- T is the subsets created from splitting set S by attribute A such that $S = \cup_{t \in T} t$
- $p(t)$ is the proportion of the number of elements in t to the number of elements in set S
- $H(t)$ is the entropy of subset t .

2.2 C4.5 and CART

C4.5 is the successor to ID3 and removed the restriction that features must be categorical by dynamically defining a discrete attribute (based on numerical variables) that partitions the continuous attribute value into a discrete set of intervals. C4.5 converts the trained trees (i.e. the output of the ID3 algorithm) into sets of if-then rules. These accuracy of each rule is then evaluated to determine the order in which they should be applied. Pruning is done by removing a rules precondition if the accuracy of the rule improves without it.

C5.0 is Quinlans latest version release under a proprietary license. It uses less memory and builds smaller rulesets than C4.5 while being more accurate.

CART (Classification and Regression Trees) is very similar to C4.5, but it differs in that it supports numerical target variables (regression) and does not compute rule sets. CART constructs binary trees using the feature and threshold that yield the largest information gain at each node.

3 Tasks

- Given the training dataset `adult.data` and the testing dataset `adult.test`, please accomplish the prediction task to determine whether a person makes over 50K a year in `adult.test` by using ID3 (or C4.5, CART) algorithm (C++ or Python), and compute the accuracy.
 1. You can process the continuous data with **bi-partition** method.
 2. You can use prepruning or postpruning to avoid the overfitting problem.
 3. You can assign probability weights to solve the missing attributes (data) problem.
- Please finish the experimental report named `E10_YourNumber.pdf`, and send it to `ai_2020@foxmail.com`

4 Codes and Results

Code:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import pickle as pk
5 import seaborn as sns
6
7 header = [ 'age', 'workclass', 'fnlwgt', 'education', 'education-num',
8            'marital-status', 'occupation', 'relationship', 'race', 'sex',
9            'capital-gain', 'capital-loss', 'hours-per-week', 'native-
            country', 'Salaries' ]
10 train_data_path = 'adult.data'
11 test_data_path = 'adult.test'
12 train_data = pd.read_csv(train_data_path, names=header)
13 test_data = pd.read_csv(test_data_path, names=header)
14 test_data.drop(0, inplace=True)
15 test_data.reset_index(drop=True, inplace=True)
16 train_data.replace('?', np.nan, inplace=True)
17 train_data.fillna(train_data.mode().iloc[0], inplace=True)
18 test_data.replace('?', np.nan, inplace=True)
19 test_data.fillna(test_data.mode().iloc[0], inplace=True)
20 continuous_cols = [ 'age', 'fnlwgt', 'education-num', 'capital-gain', '
    capital-loss', 'hours-per-week' ]
21 pos1 = int(len(train_data)/3)
22 pos2 = 2 * pos1
23 intervals = {}
24
25 for col in continuous_cols:
26     i1 = sorted(train_data[col])[pos1]
27     i2 = sorted(train_data[col])[pos2]
28     intervals[col] = (range(0, i1+1), range(i1+1, i2+1), range(i2+1,
        sorted(train_data[col])[len(train_data)-1]+1))
29
30 rev_intervals = {}
31 for k, v in intervals.items():
```

```

32 tmp = {}
33 for idx, r in enumerate(v):
34     for i in r:
35         tmp[i] = idx
36 rev_intervals[k] = tmp
37
38 dsp_dict = {
39     1: [ 'Private', 'Self-emp-not-inc', 'Self-emp-inc', 'Federal-gov', '
        Local-gov', 'State-gov', 'Without-pay', 'Never-worked' ],
40     3: [ 'Bachelors', 'Some-college', '11th', 'HS-grad', 'Prof-school', '
        Assoc-acdm', 'Assoc-voc', '9th', '7th-8th', '12th', 'Masters', '1
        st-4th', '10th', 'Doctorate', '5th-6th', 'Preschool' ],
41     5: [ 'Married-civ-spouse', 'Divorced', 'Never-married', 'Separated',
        'Widowed', 'Married-spouse-absent', 'Married-AF-spouse' ],
42     6: [ 'Tech-support', 'Craft-repair', 'Other-service', 'Sales', 'Exec-
        managerial', 'Prof-specialty', 'Handlers-cleaners', 'Machine-op-
        inspct', 'Adm-clerical', 'Farming-fishing', 'Transport-moving', '
        Priv-house-serv', 'Protective-serv', 'Armed-Forces' ],
43     7: [ 'Wife', 'Own-child', 'Husband', 'Not-in-family', 'Other-relative',
        'Unmarried' ],
44     8: [ 'White', 'Asian-Pac-Islander', 'Amer-Indian-Eskimo', 'Other', '
        Black' ],
45     9: [ 'Female', 'Male' ],
46     13: [ 'United-States', 'Cambodia', 'England', 'Puerto-Rico', 'Canada',
        'Germany', 'Outlying-US(Guam-USVI-etc)', 'India', 'Japan', '
        Greece', 'South', 'China', 'Cuba', 'Iran', 'Honduras', '
        Philippines', 'Italy', 'Poland', 'Jamaica', 'Vietnam', 'Mexico',
        'Portugal', 'Ireland', 'France', 'Dominican-Republic', 'Laos', '
        Ecuador', 'Taiwan', 'Haiti', 'Columbia', 'Hungary', 'Guatemala',
        'Nicaragua', 'Scotland', 'Thailand', 'Yugoslavia', 'El-Salvador',
        'Trinidad&Tobago', 'Peru', 'Hong', 'Holand-Netherlands' ]
47 }
48
49 def dsp2numlist(idx):
50     return list(range(len(dsp_dict[idx])))
51
52 AttrSet = [
53     (0, [0, 1, 2], 'age'),
54     (1, dsp2numlist(1), 'workclass'),
55     (2, [0, 1, 2], 'fnlwgt'),
56     (3, dsp2numlist(3), 'education'),
57     (4, [0, 1, 2], 'education-num'),
58     (5, dsp2numlist(5), 'marital-status'),
59     (6, dsp2numlist(6), 'occupation'),
60     (7, dsp2numlist(7), 'relationship'),
61     (8, dsp2numlist(8), 'race'),
62     (9, dsp2numlist(9), 'sex'),
63     (10, [0, 1, 2], 'capital-gain'),
64     (11, [0, 1, 2], 'capital-loss'),

```

```

65     (12, [0, 1, 2], 'hours-per-week'),
66     (13, dsp2numlist(13), 'native-country')
67 ]
68
69 train_label = [1 if val == '>50K' else 0 for val in train_data['
    Salaries']]
70 train_input = []
71
72 for idx in range(len(train_data)):
73     tmp = [dsp_dict[i].index(val.strip()) if int(i) in dsp_dict.keys()
74            else rev_intervals[train_data.columns[i]].get(val, 2) for i,
75            val in enumerate(train_data.iloc[idx][: -1])]
76     train_input.append(tmp)
77 test_label = [1 if val == '>50K.' else 0 for val in test_data['Salaries
    ']]
78 test_input = []
79 for idx in range(len(test_data)):
80     tmp = [dsp_dict[i].index(val.strip()) if int(i) in dsp_dict.keys()
81            else rev_intervals[test_data.columns[i]].get(val, 2) for i,
82            val in enumerate(test_data.iloc[idx][: -1])]
83     test_input.append(tmp)
84
85 def Entropy(Data):
86     labels = [sample[-1] for sample in Data]
87     types = set(labels)
88     types_counts = [labels.count(type) for type in types]
89     probs = [prob/len(Data) for prob in types_counts]
90     return -np.sum(probs*np.log2(probs))
91
92 def Gain(Data, attr):
93     entropy = Entropy(Data)
94     attr_num = attr[0]
95     attr_vals = attr[1]
96     entropys = [0 for val in attr_vals]
97     weights = [0 for val in attr_vals]
98     for idx, val in enumerate(attr_vals):
99         sub_data = []
100         for sample in Data:
101             if sample[attr_num] == val:
102                 sub_data.append(sample)
103                 weights[idx] += 1
104             entropys[idx] = Entropy(sub_data)
105             weights[idx] /= len(Data)
106     return entropy - np.sum(np.multiply(weights, entropys))
107
108 def Gini(Data):
109     labels = [sample[-1] for sample in Data]
110     types = set(labels)
111     types_counts = [labels.count(type) for type in types]

```

```

110     probs = [prob/len(Data) for prob in types_counts]
111     return 1 - np.sum(np.power(probs, 2))
112
113 def Gini_index(Data, attr):
114     gini = Gini(Data)
115     attr_num = attr[0]
116     attr_vals = attr[1]
117     ginis = [0 for val in attr_vals]
118     weights = [0 for val in attr_vals]
119     for idx, val in enumerate(attr_vals):
120         sub_data = []
121         for sample in Data:
122             if sample[attr_num] == val:
123                 sub_data.append(sample)
124                 weights[idx] += 1
125         ginis[idx] = Gini(sub_data)
126         weights[idx] /= len(Data)
127     return np.sum(np.multiply(weights, ginis))
128
129 def chooseBestAttr(Data, Attrset, method='ID3'):
130     best_attr = Attrset[0]
131     best_gain = -1
132     best_gini = np.Inf
133     for attr_tuple in Attrset:
134         gain = Gain(Data, attr_tuple)
135         best_attr = attr_tuple if gain > best_gain else best_attr
136         best_gain = gain if gain > best_gain else best_gain
137     return best_attr
138
139 def splitData(Data, attr_num, attr_val):
140     sub_data = []
141     for sample in Data:
142         if sample[attr_num] == attr_val:
143             sub_data.append(sample[:attr_num] + sample[attr_num+1:])
144     return sub_data
145
146 def getMajority(Data):
147     labels = [sample[-1] for sample in Data]
148     types = list(set(labels))
149     types_counts = [labels.count(type) for type in types]
150     major = 0
151     max_count = 0
152     for idx, type_count in enumerate(types_counts):
153         major = types[idx] if max_count < type_count else major
154         max_count = type_count if max_count < type_count else max_count
155     return str(major)
156
157 def GenerateTree(Data, Attrset, method='ID3'):
158     labels = [sample[-1] for sample in Data]

```

```

159     if len(set(labels)) == 1:
160         return str(labels[0])
161     if len(Attrset) == 0:
162         return getMajority(Data)
163     flag = False
164     for attr_tuple in Attrset:
165         if len(set([sample[attr_tuple[0]] for sample in Data])) != 1:
166             flag = True
167             break
168     if not flag:
169         return getMajority(Data)
170     best_attr = chooseBestAttr(Data, Attrset, method)
171     attr_num = best_attr[0]
172     attr_vals = best_attr[1]
173     attr_name = best_attr[2]
174     for idx, attr in enumerate(Attrset):
175         if attr[0] > attr_num:
176             Attrset[idx] = (attr[0]-1, attr[1], attr[2])
177     del(Attrset[Attrset.index(best_attr)])
178     Node = {attr_name: {}}
179     for val in attr_vals:
180         sub_data = splitData(Data, attr_num, val)
181         if len(sub_data) == 0:
182             return getMajority(Data)
183         else:
184             Node[attr_name][val] = GenerateTree(sub_data, Attrset[:],
185                                                 method)
186     return Node
187
188 def Classifier(DecisionTree, AttrSet, SampleData):
189     root = list(DecisionTree.keys())[0]
190     for attr_tuple in AttrSet:
191         if root == attr_tuple[2]:
192             key = SampleData[attr_tuple[0]]
193             succ = DecisionTree[root][key]
194             if isinstance(succ, dict):
195                 return Classifier(succ, AttrSet, SampleData)
196             else:
197                 return succ
198
199 def show_accuracy(DecisionTree, AttrSet, testing_data, log=False):
200     labels = [sample[-1] for sample in testing_data]
201     res = []
202     for sample in testing_data:
203         res.append(Classifier(DecisionTree, AttrSet, sample[:-1]))
204     check = [labels[idx] + int(res[idx]) for idx in range(len(
205         testing_data))]
206     if log:
207         print("Total Accuracy: %.5f" % (1-check.count(1)/len(

```



```

        testing_data)))
206     else:
207         return 1 - check.count(1)/len(testing_data)
208
209 testing_attrset = [(0, [0, 1, 2, 3], '2nd'), (1, [0, 1, 2, 3], '3rd'),
210                  (2, [0, 1, 2, 3, 4, 5, 6], '4th')]
211 X_train = [data + [train_label[idx]] for idx, data in enumerate(
212             train_input)]
213 X_test = [data + [test_label[idx]] for idx, data in enumerate(test_input
214                                                                )]
215 SalaryPredict_DT_ID3 = GenerateTree(X_train, AttrSet[:])
216 print("=="*10+ ' Testing ID3 '+'=="*10)
217 show_accuracy(SalaryPredict_DT_ID3, AttrSet, X_test, True)

```

Result:

```

heze@HeZes-MBP > python DT.py
==== Testing ID3 =====
Total Accuracy: 0.82212

```